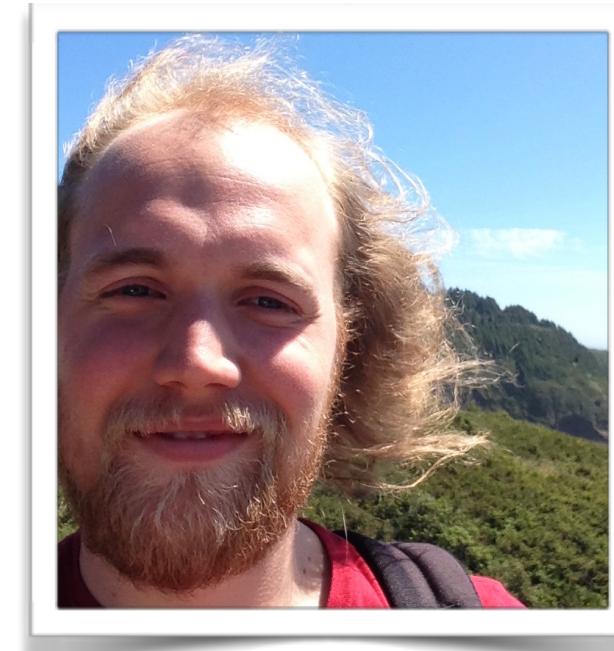


C_b: A New Modular Approach to Implementing Efficient and Tunable Collections

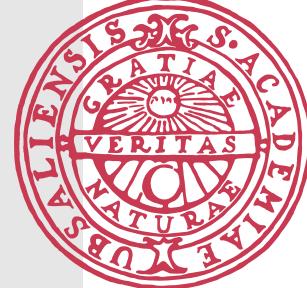
Stephan Brandauer
Uppsala University



Elias Castegren
KTH

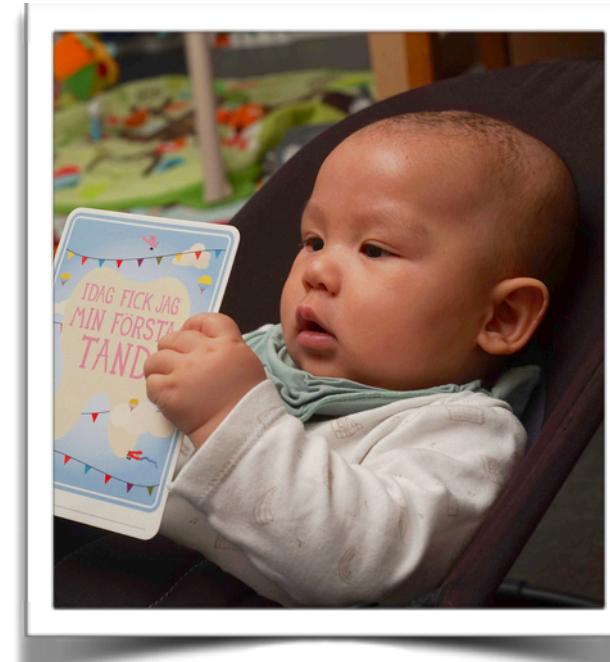


Tobias Wrigstad
Uppsala University

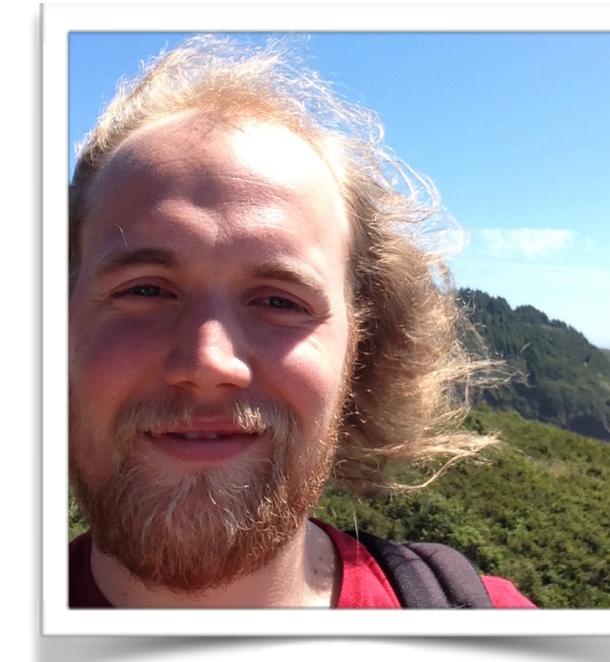


Cb: A New Modular Approach to Implementing Efficient and Tunable Collections

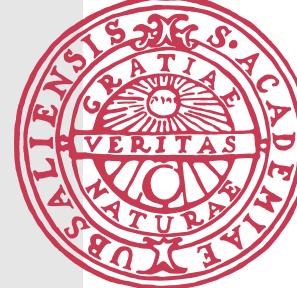
Stephan Brandauer
Uppsala University



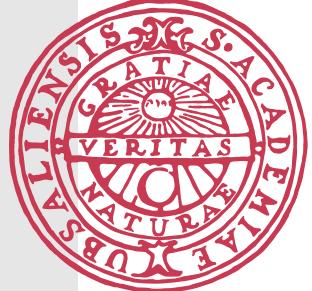
Elias Castegren
KTH



Tobias Wrigstad
Uppsala University

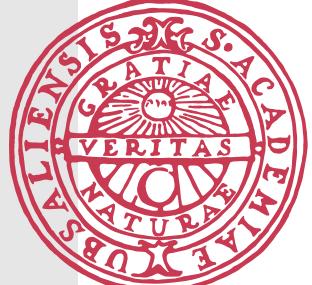


Motivation



Motivation

- In the implementation of collections, the business logic and the data representation are intimately tangled



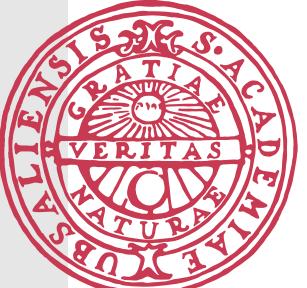
Motivation

- In the implementation of collections, the business logic and the data representation are intimately tangled
- Two implementations of the same interface are functionally equivalent but may differ on non-functional properties

Prepend of a linked list is fast — prepend of an array list is not

Random access in an array list is fast — random access in a linked list is not

...



Motivation

- In the implementation of collections, the business logic and the data representation are intimately tangled
- Two implementations of the same interface are functionally equivalent but may differ on non-functional properties

Prepend of a linked list is fast — prepend of an array list is not

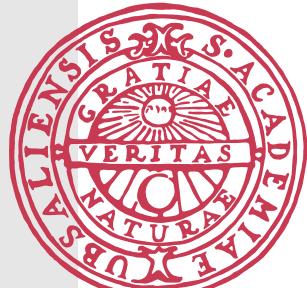
Random access in an array list is fast — random access in a linked list is not

- In Java, for example, we abstract the concrete implementation (`LinkedList`, `ArrayList`) by programming against a `List` interface

As we learn about the software we are developing, we can easily swap out concrete implementations

Also possible "just-in-time", transitioning from prepend-phase to random access-phase, etc.

...



Implementation of LinkedList and ArrayList

- They both implement the List interface with **41** methods
- LinkedList implemented by Josh Bloch, ArrayList implemented by Josh Bloch and Neal Gafter
- They share ~**5** methods (out of ~**50**) through inheritance from e.g., AbstractList

Examples: equals, hashCode, toString, containsAll

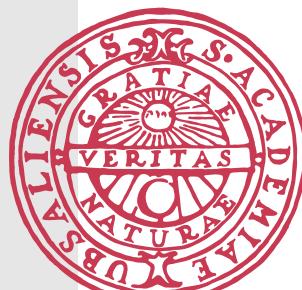
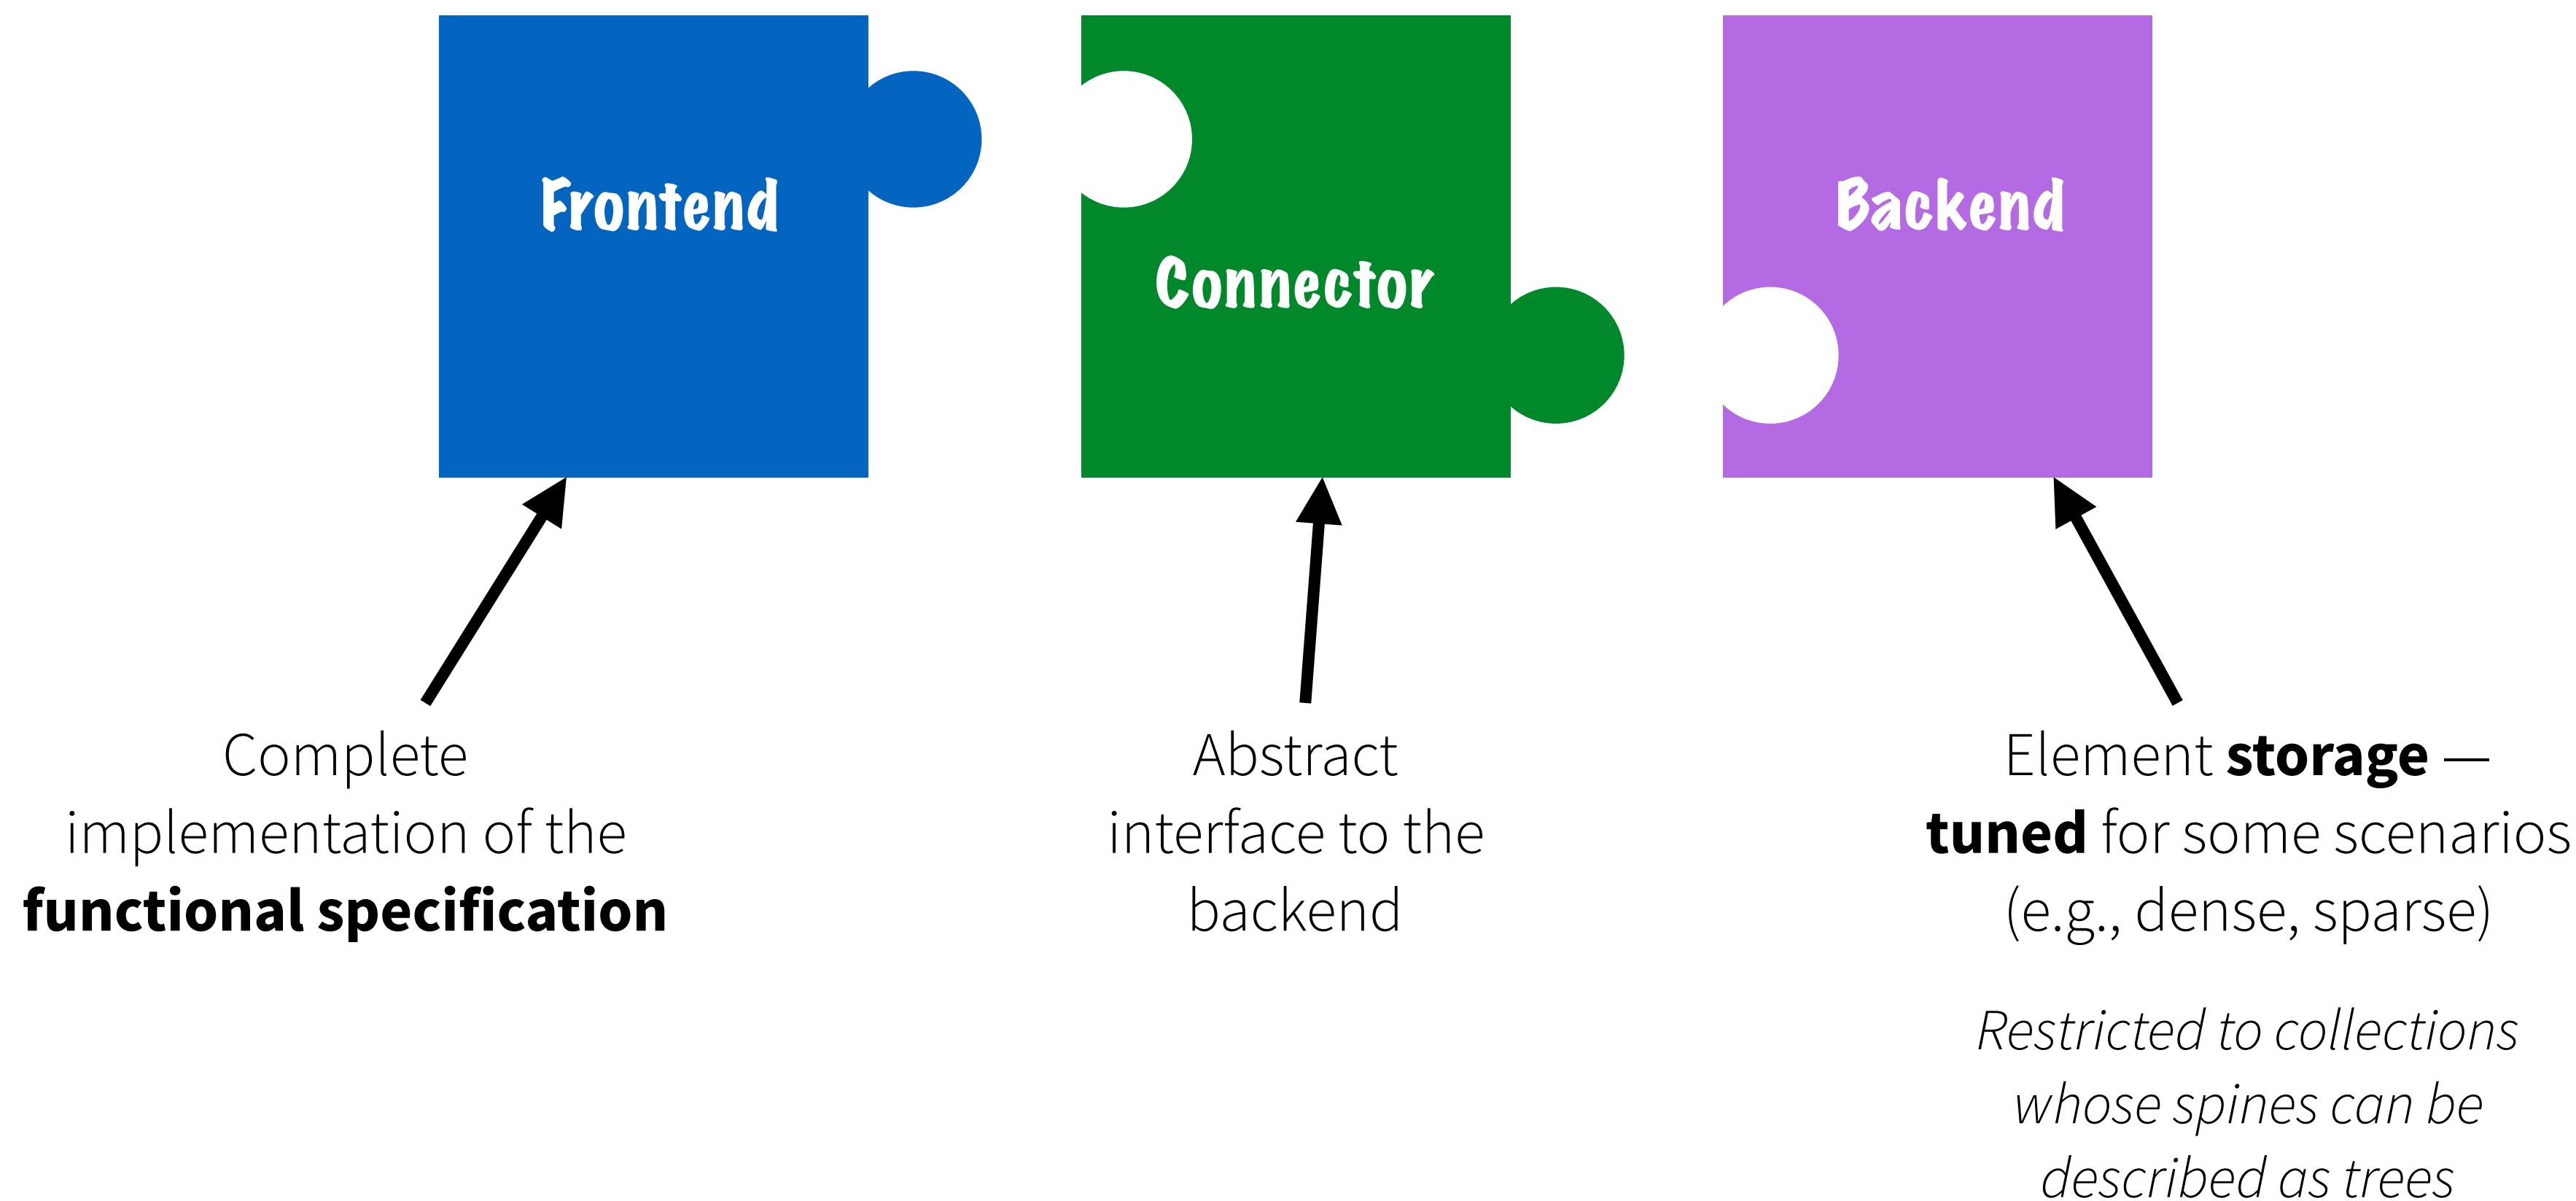
- In the source code, there are more opportunities for reuse, but mostly "ignored" for performance reasons because different data representation favor different access patterns



Numbers included to show magnitude

This Work is About the Separation of Data Representation from the Implementation of Functional Specifications

- Collections are implemented as three components working together:



This Work is About the Separation of Data Representation from the Implementation of Functional Specifications

- Collections are implemented as three components working together:

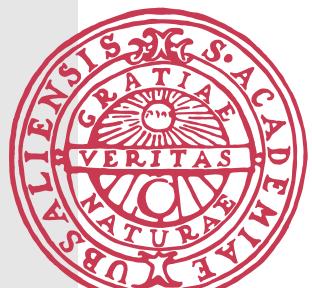


- Frontends and backends can be combined freely

List interface + Array-based backend = ArrayList

Connector logic synthesized from an abstract specification

- Prototype Implementation in Java



This Work is About the Separation of Data Representation from the Implementation of Functional Specifications

- Collections are implemented as three components working together:

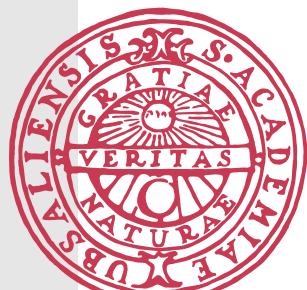


- Frontends and backends can be combined freely

List interface + Array-based backend = ArrayList

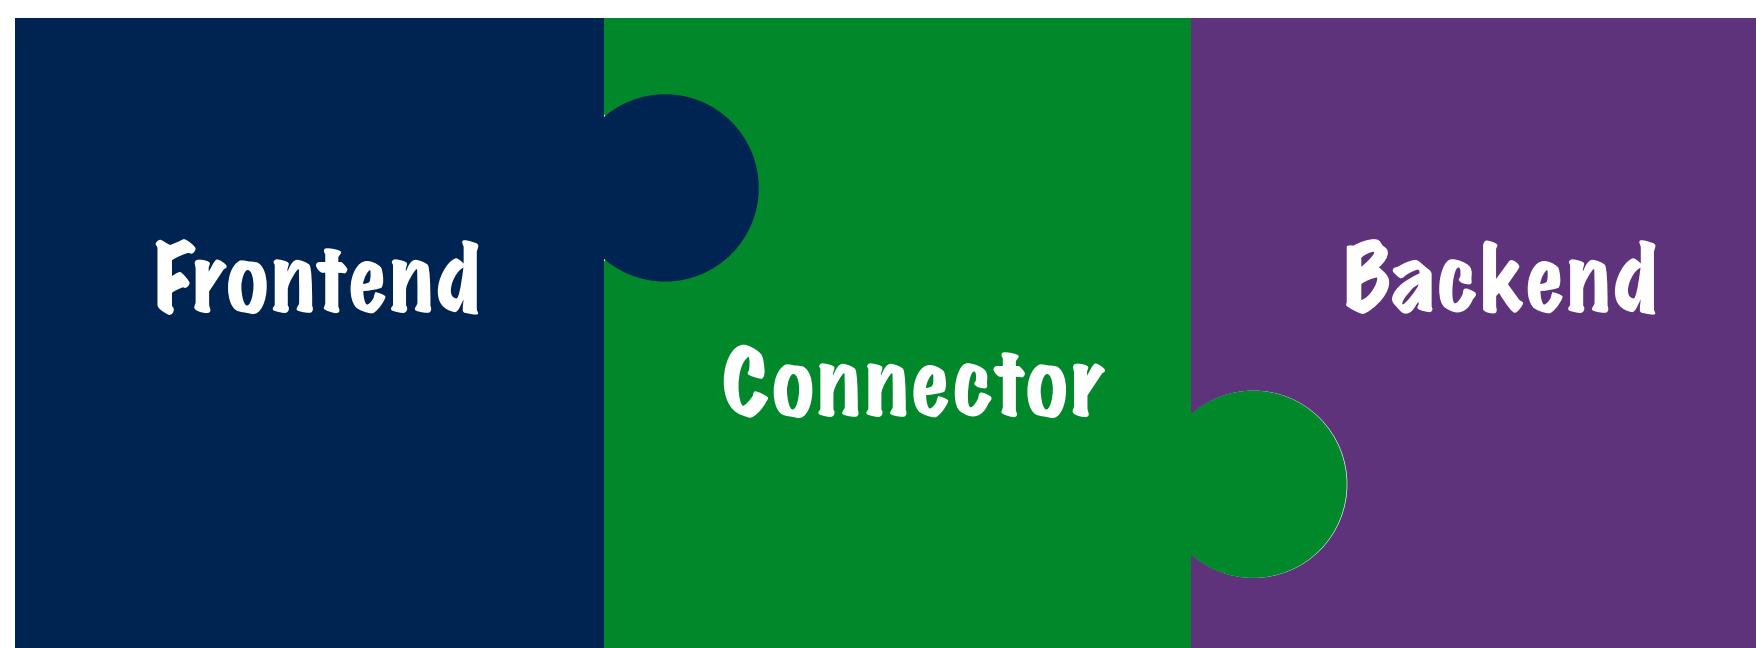
Connector logic synthesized from an abstract specification

- Prototype Implementation in Java



This Work is About the Separation of Data Representation from the Implementation of Functional Specifications

- Collections are implemented as three components working together:

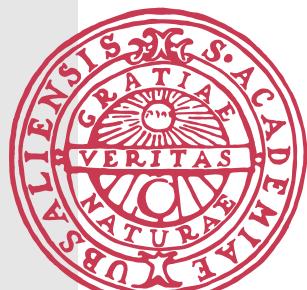


- Frontends and backends can be combined freely

List interface + Array-based backend = ArrayList

Connector logic synthesized from an abstract specification

- Prototype Implementation in Java



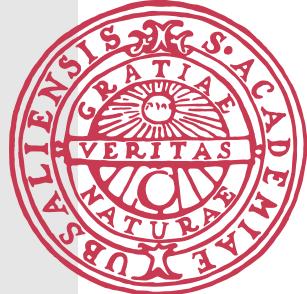
Connector Logic

- Front-ends access storage via a cursor, which can be thought of as a pointer variable with strict rules on how it may be reassigned

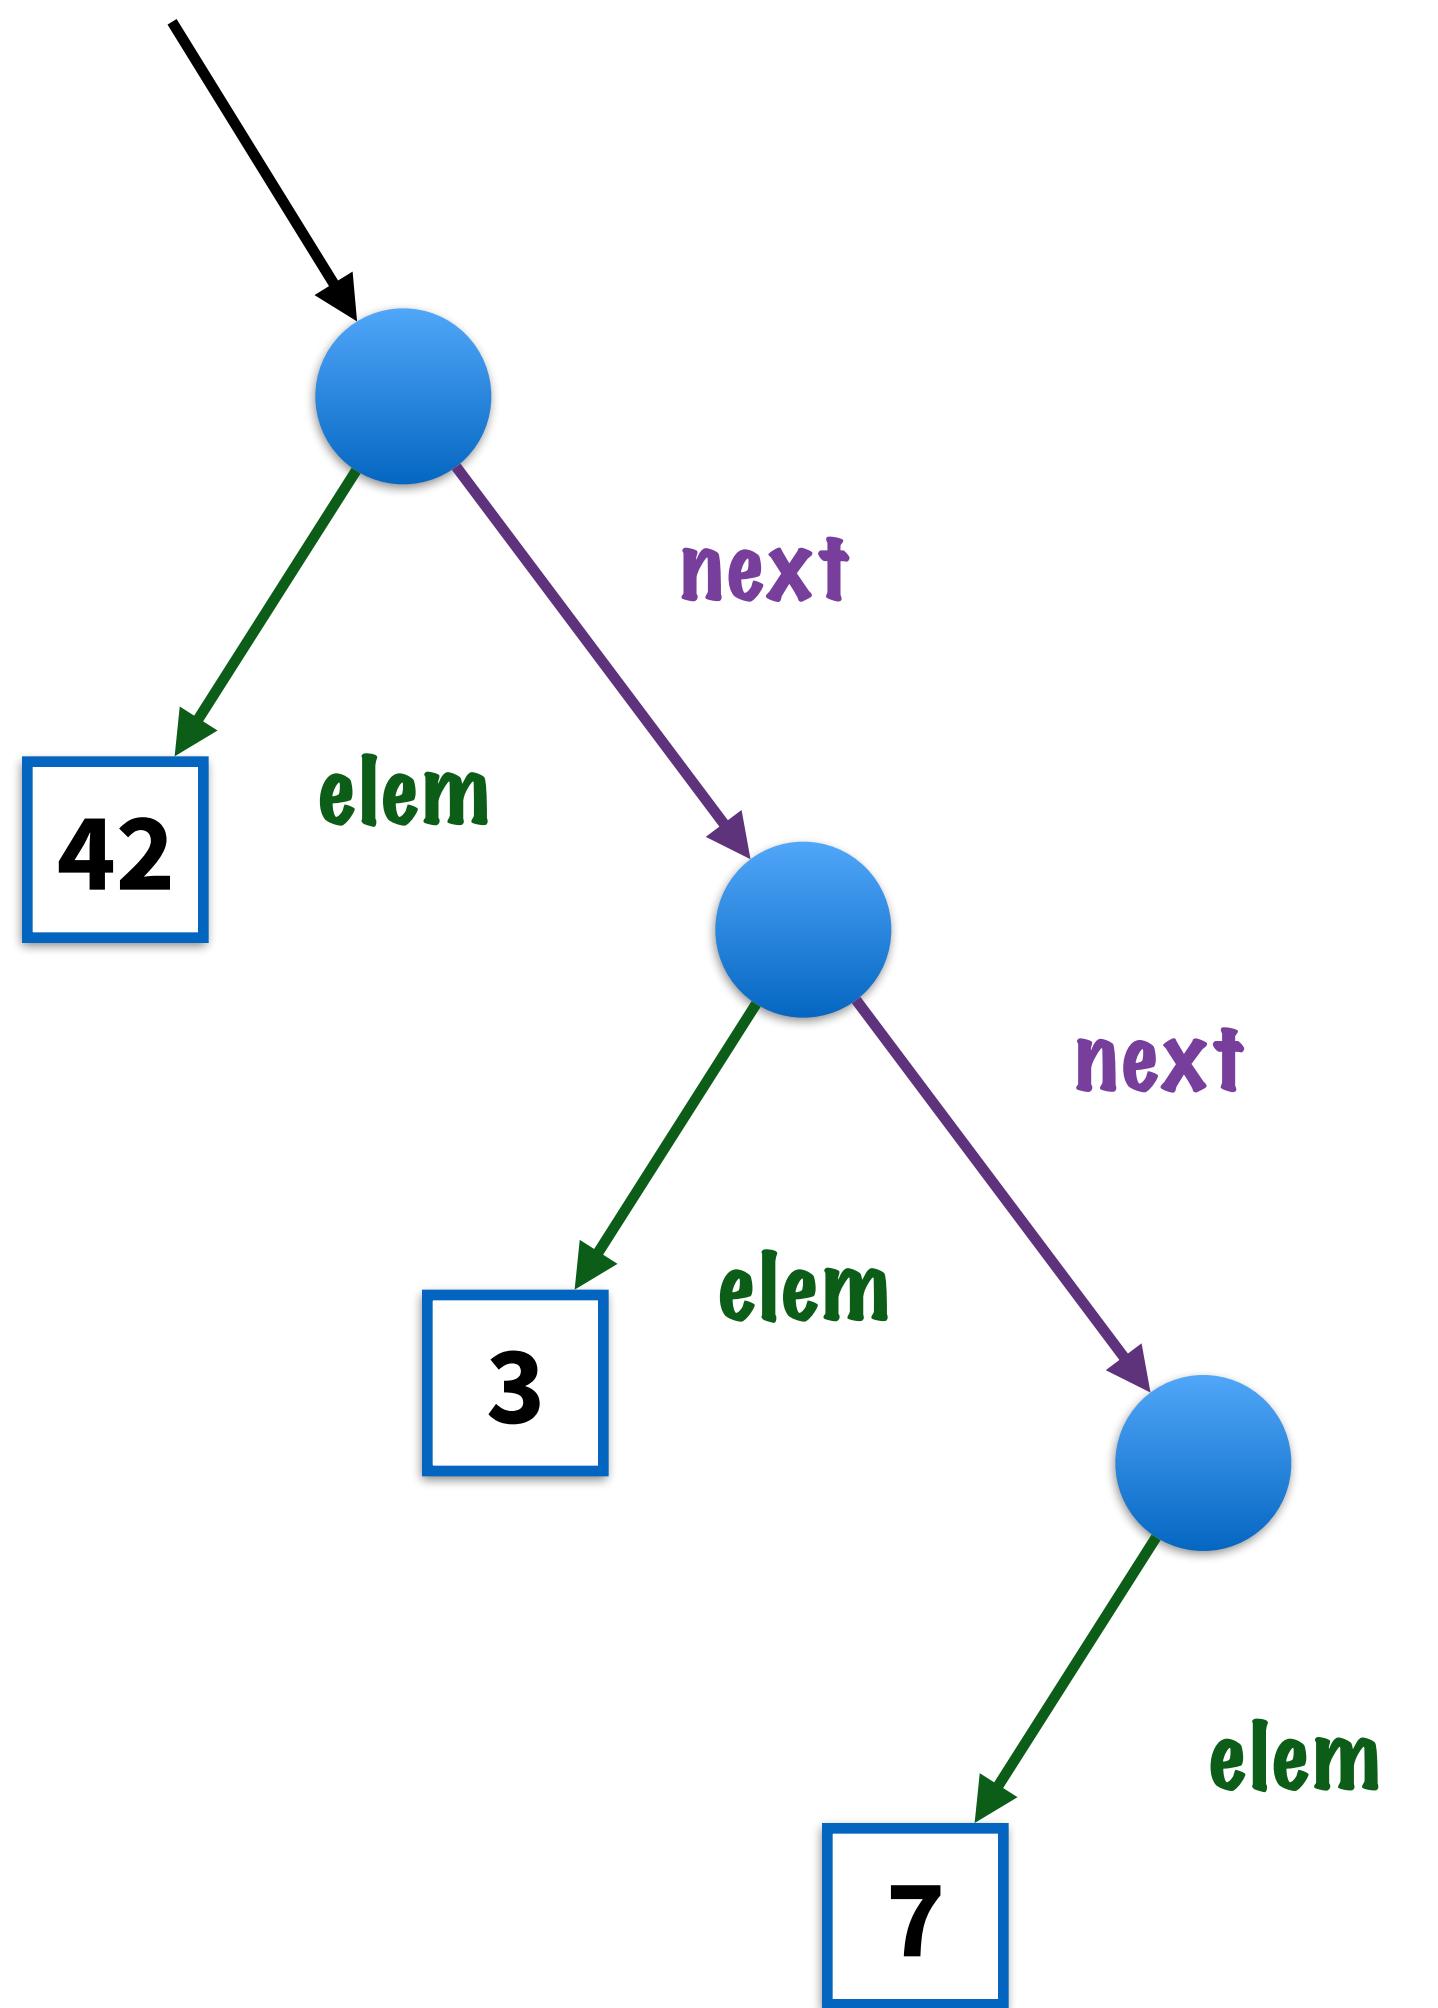
The cursor is a robot that navigates the storage in the backend and retrieves or stores elements on their uniquely named storage locations

How the cursor may move is specified as an automaton

This automaton is embedded in the source code as a Java annotation, cursor generated from it automatically

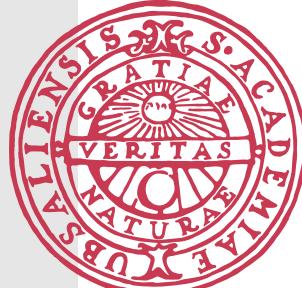


Logical representation of a sequence/list



Connector automaton:

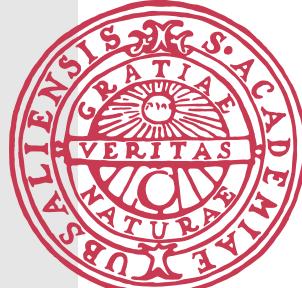
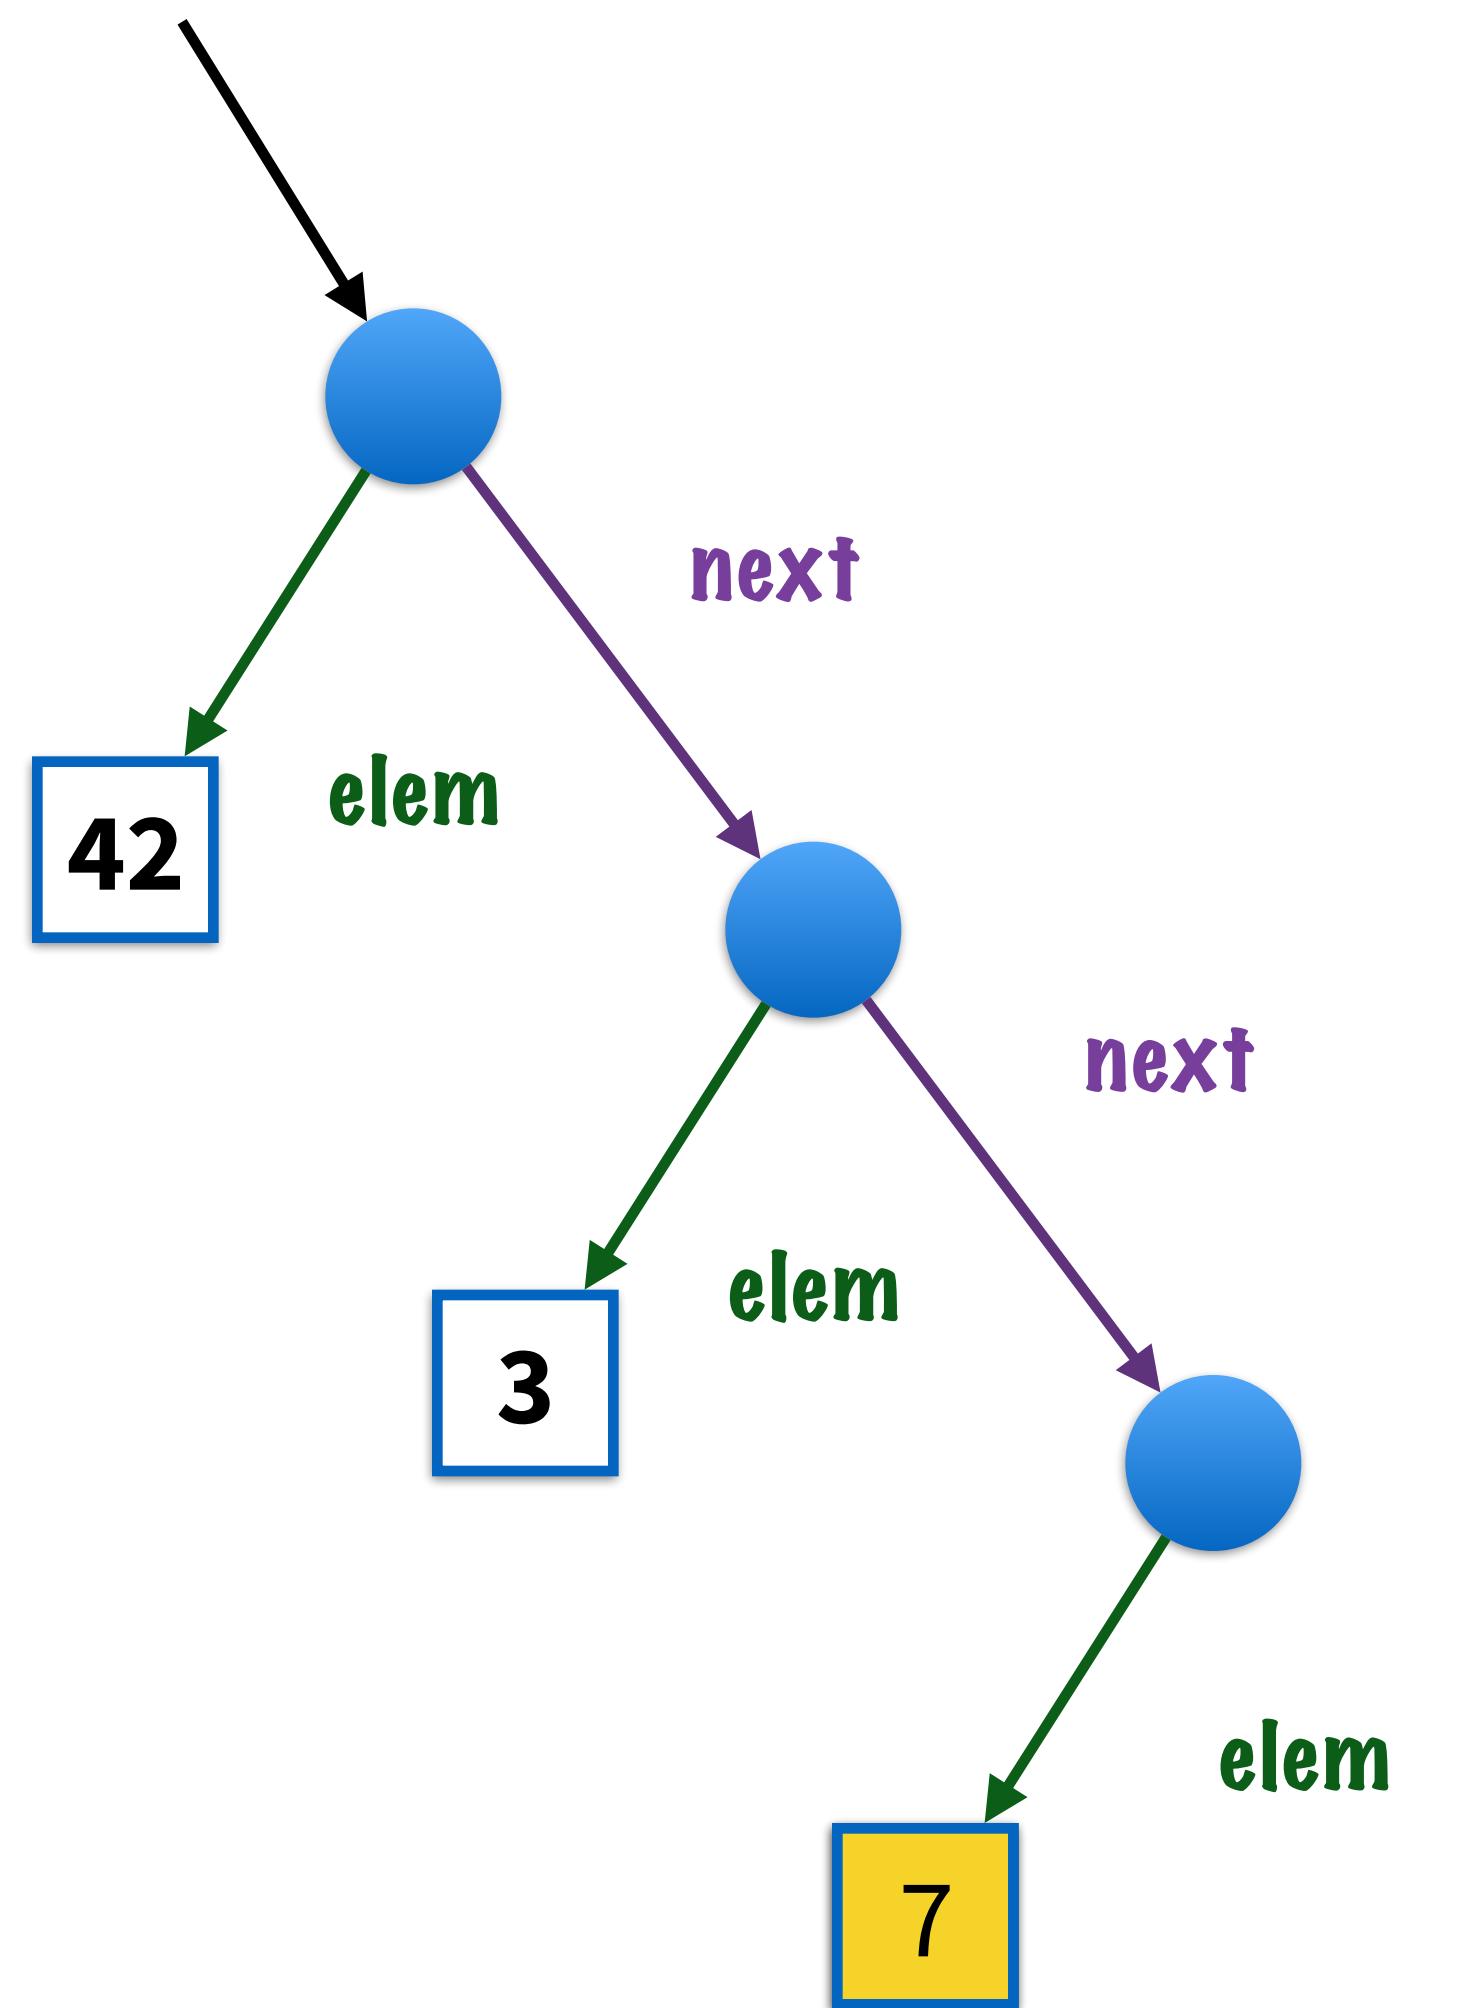
$*(\text{next}) \rightarrow \text{elem}$

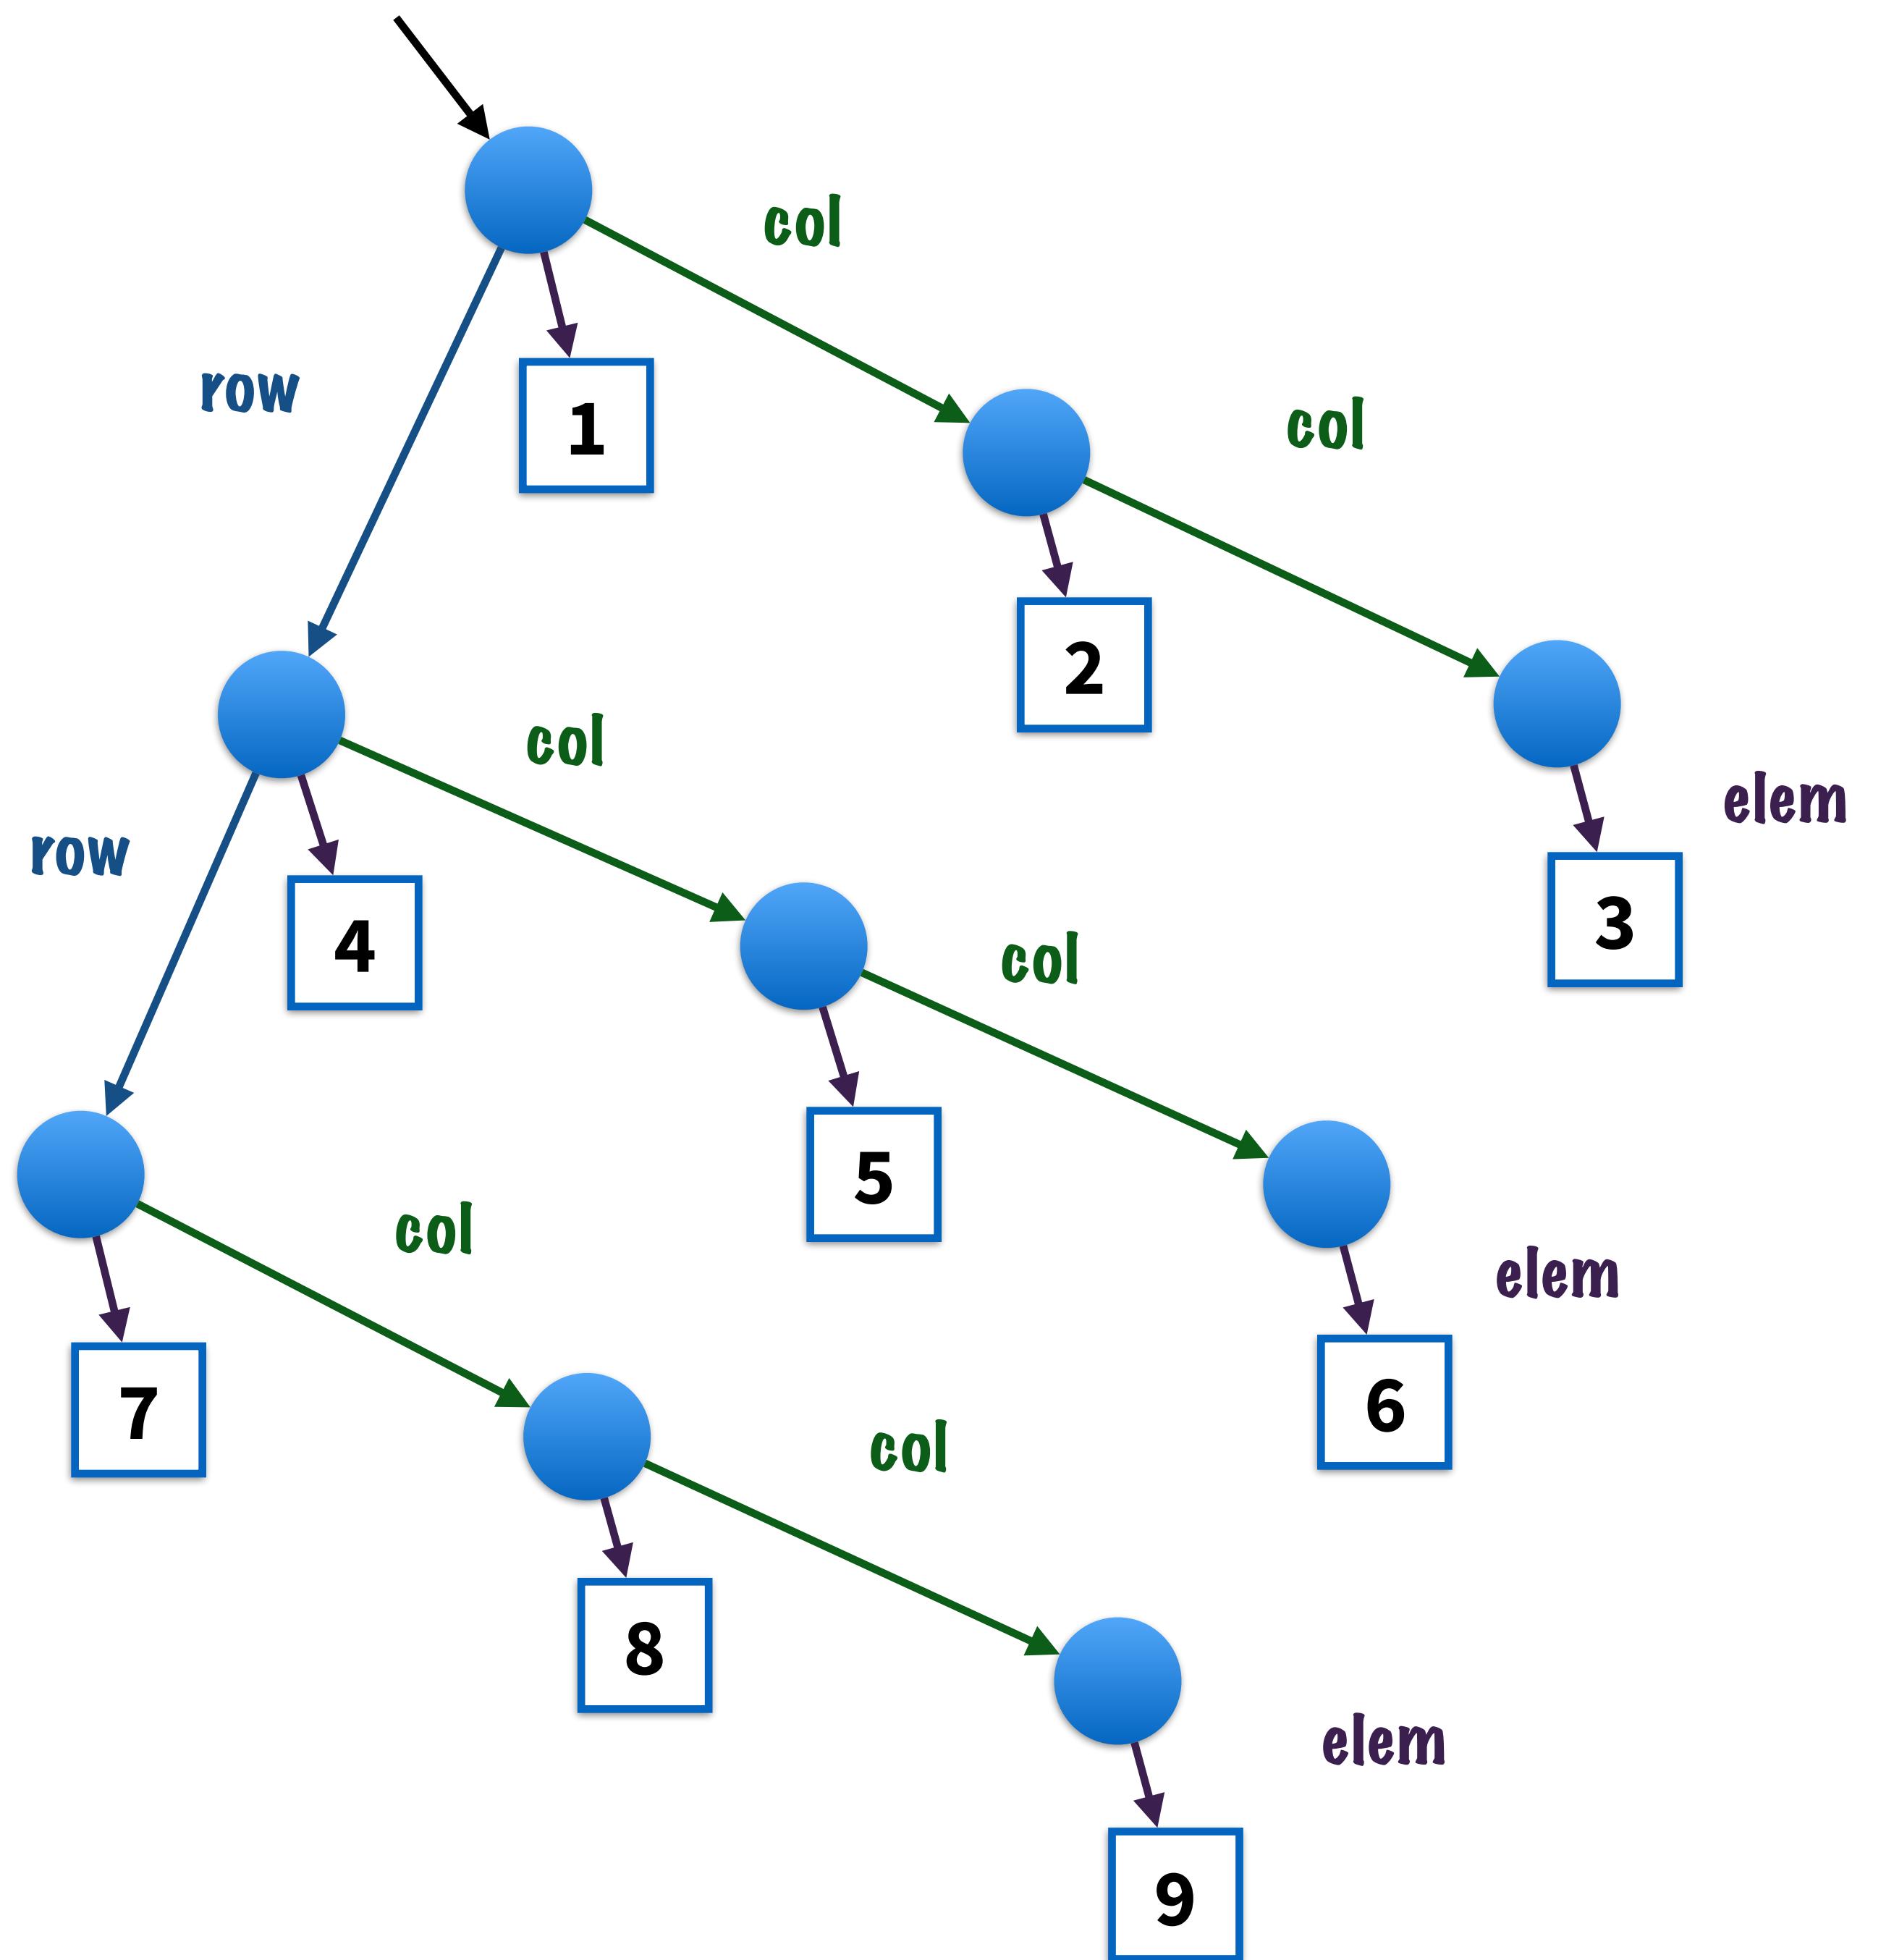


Logical representation of a sequence/list

Connector automaton:

${}^*(\text{next}) \rightarrow \text{elem}$

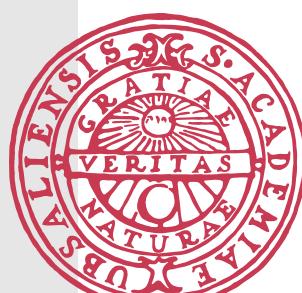


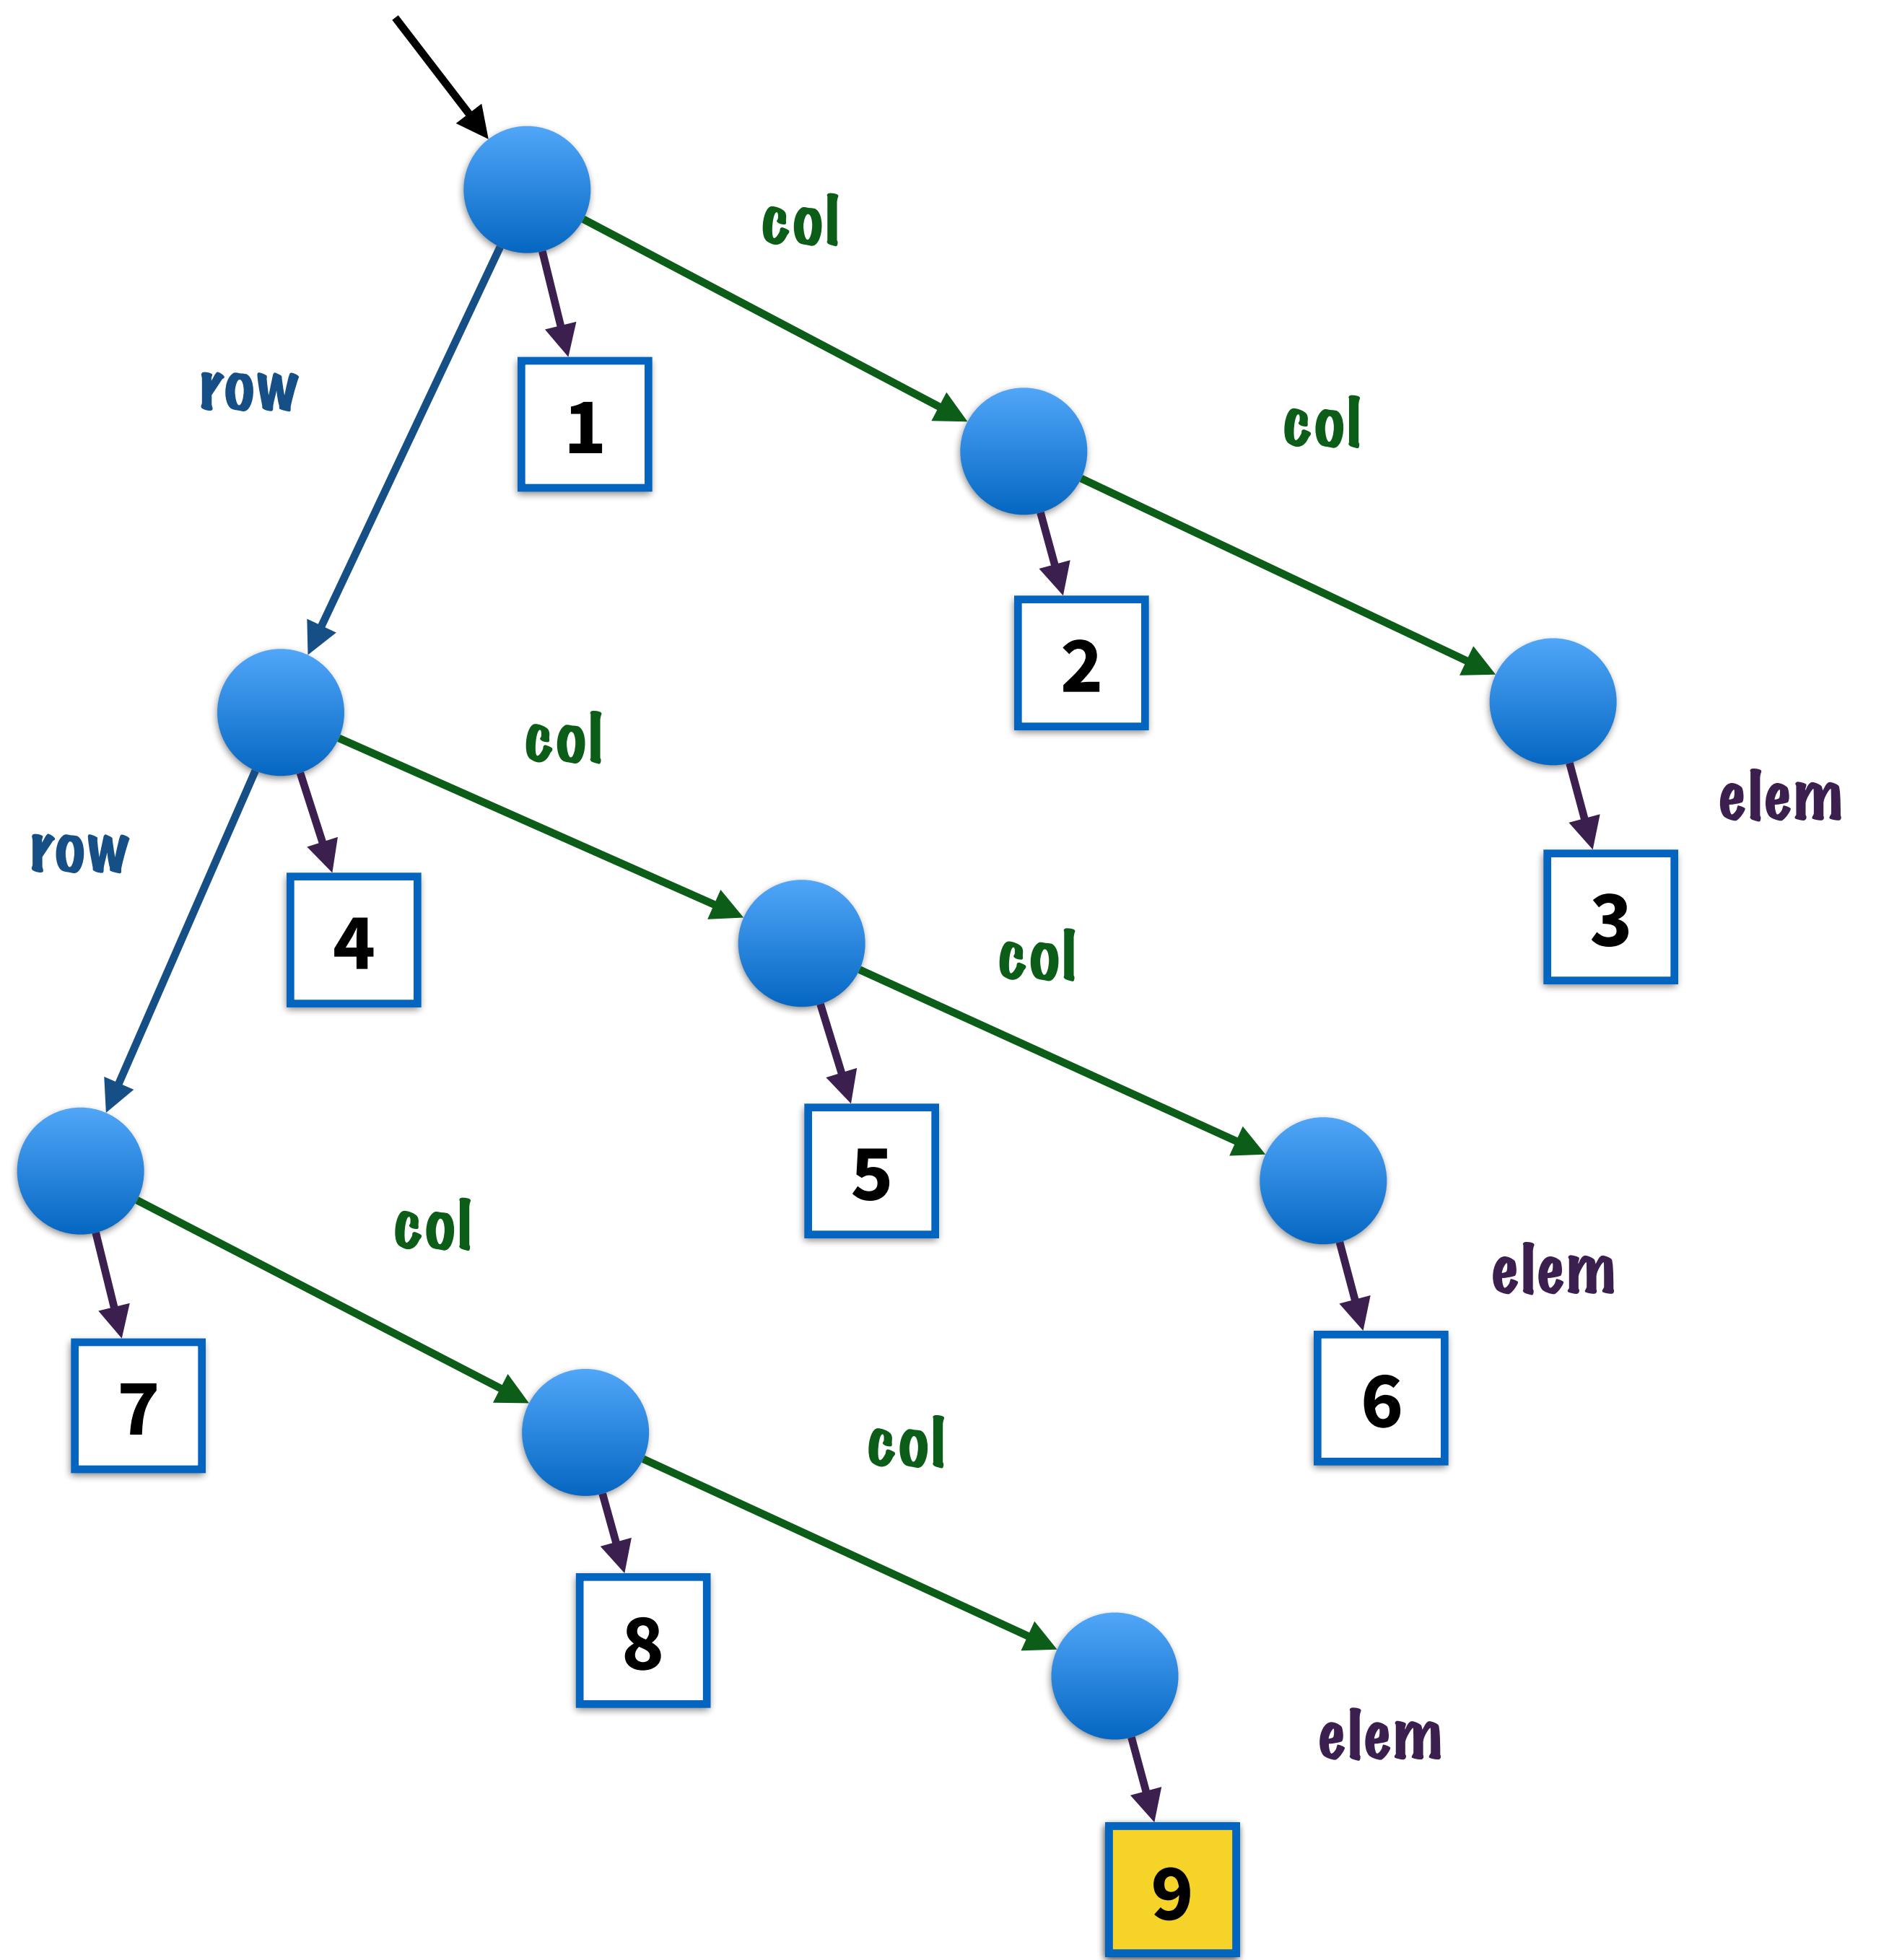


Logical representation of a
 3×3 matrix

Connector automaton:

$*3(\text{row}) \rightarrow *3(\text{col}) \rightarrow \text{elem}$

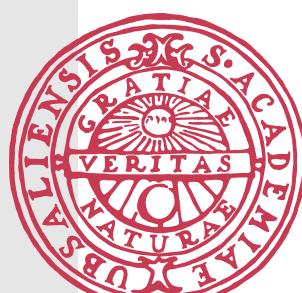




Logical representation of a
 3×3 matrix

Connector automaton:

$*3(\text{row}) \rightarrow *3(\text{col}) \rightarrow \text{elem}$

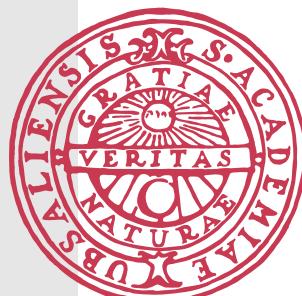


~~Legal Robot Movement~~ Cursor Automaton

C _b expression	Interpretation	Example
$*(\text{next}) \rightarrow \text{elem}$	Sequence	$\text{next}.\text{next}.\text{next}.\text{elem}$
$*(\ (\text{left}, \text{right})) \rightarrow \text{elem}$	2-ary tree	$\text{left}.\text{right}.\text{left}.\text{elem}$
$*(*7(\text{child}) \rightarrow \text{pick}) \rightarrow \text{elem}$	7-ary tree	$\text{child}.\text{child}.\text{pick}.\text{child}.\text{pick}.\text{elem}$
$*3(\text{row}) \rightarrow *3(\text{col}) \rightarrow \text{elem}$	3 x 3 matrix	$\text{row}.\text{row}.\text{col}.\text{col}.\text{elem}$

$C ::=$ *(C_b expression.)*

- | $*(S) \rightarrow S'$ Repeat S unboundedly, then S' .
- | $*N(S) \rightarrow S'$ Repeat S upto N times, then S' .
- | $\|(A_1, \dots, A_n)$ Pick one alternative in $\{A_1, \dots, A_n\}$.



Backend Abstraction

- Each storage location has a unique path to it
- We can thus enumerate the paths and use these as ids

For example, `row().row().col().col().elem()` on a previous slide will have id 8 – which nicely translates into an array index, or key

The equivalent path `row(2).col(2).elem()` can be implemented in constant-time for backends that support it

- The key components of a backend are `get()` and `set()` operations that take legal paths as inputs

```
// Backend
Backend<T> backend = ...

// Create a new cursor, positioned at root
Cursor c = new Cursor();

// Move the cursor to the right place
c.row().row().col().col();

// Extract the element
T value = backend.get(c.elem());
```



Backend Abstraction

- Each storage location has a unique path to it
- We can thus enumerate the paths and use these as ids

For example, `row().row().col().col().elem()` on a previous slide will have id 8 – which nicely translates into an array index, or key

The equivalent path `row(2).col(2).elem()` can be implemented in constant-time for backends that support it

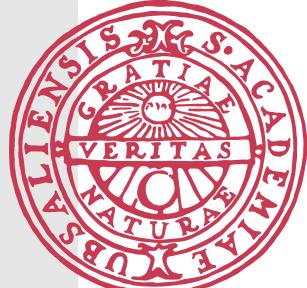
- The key components of a backend are `get()` and `set()` operations that take legal paths as inputs

```
// Backend
Backend<T> backend = ...

// Create a new cursor, positioned at root
Cursor c = new Cursor();

// Analogous to matrix[2][2] but abstract
c.row(2).col(2);

// Extract the element
T value = backend.get(c.elem());
```



Backend Abstraction

- Each storage location has a unique path to it
- We can thus enumerate the paths and use these as ids

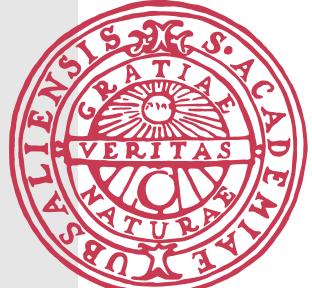
For example, `row().row().col().col().elem()` on a previous slide will have id 8 – which nicely translates into an array index, or key

The equivalent path `row(2).col(2).elem()` can be implemented in constant-time for backends that support it

- The key components of a backend are `get()` and `set()` operations that take legal paths as inputs

This call actually returns the int 8

```
// Backend  
Backend<T> backend = ...  
  
// Create a new cursor, positioned at root  
Cursor c = new Cursor();  
  
// Analogous to matrix[2][2] but abstract  
c.row(2).col(2);  
  
// Extract the element  
T value = backend.get(c.elem());
```



Generation of Cursor Classes

- Generated at compile time by a library — no compiler modifications
- Encapsulates logic for calculating and returning the unique id for a path, represented as an integer
- Stateful — pointer variable analogy

Given methods from the symbols in the C_b expression — e.g., `left`, `right`, `next`, `elem` etc.

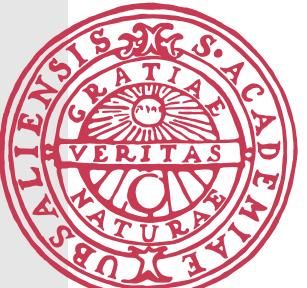
Each call to a method will manipulate the id state

`next()` = `++`

`left()` = `current position * 2 + 1` (*binary heap layout*)

`right()` = `current position * 2 + 2`

`elem()` returns the path identifier

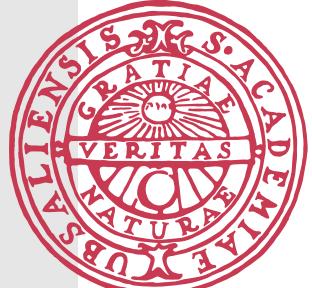


Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();

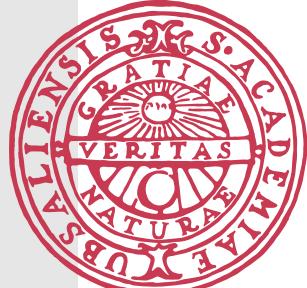
    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```



Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem") ←  
class Sequence<T> implements java.util.List {  
    private Backend<T> data;  
    private SequenceCursor tail = new SequenceCursor();  
  
    boolean add(T obj) {  
  
        data.set(this.tail.elem(), obj);  
  
        tail.next();  
        return true;  
    }  
  
    boolean contains(Object obj) {  
        return data.usedLocations().findAny((T) obj) != -1;  
    }  
}
```



Implementing Collections using an Abstract Representation

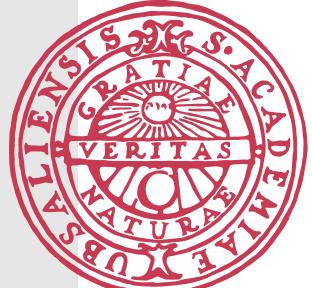
```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data; ←
    private SequenceCursor tail = new SequenceCursor();

    boolean add(T obj) {

        data.set(this.tail.elem(), obj);

        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```

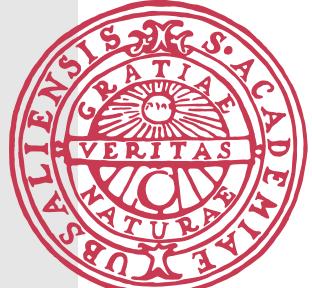


Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();
    ←

    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```

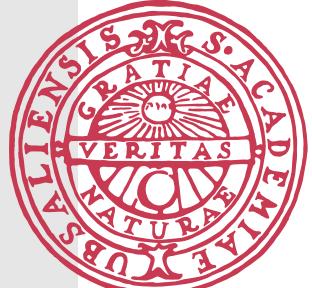
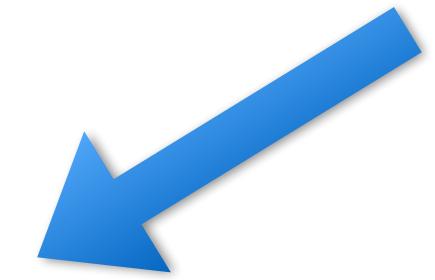


Implementing Collections using an Abstract Representation

```
@Cflat("*(next)->elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();

    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```

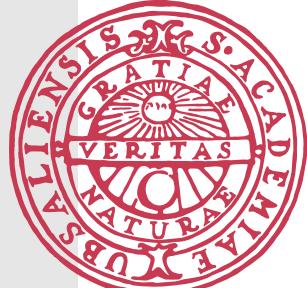


Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();

    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```

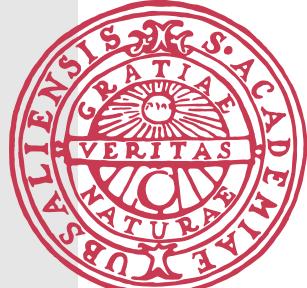


Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();

    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next(); ←
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```

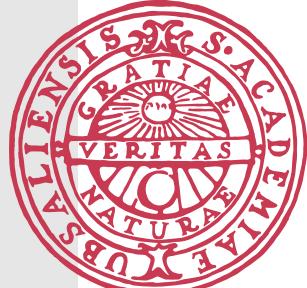


Implementing Collections using an Abstract Representation

```
@Cflat("* (next) -> elem")
class Sequence<T> implements java.util.List {
    private Backend<T> data;
    private SequenceCursor tail = new SequenceCursor();

    boolean add(T obj) {
        data.set(this.tail.elem(), obj);
        tail.next();
        return true;
    }

    boolean contains(Object obj) {
        return data.usedLocations().findAny((T) obj) != -1;
    }
}
```



Backend Implementation

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

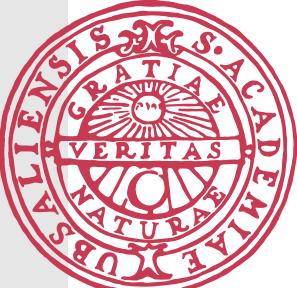
```
public interface Backend<T> {  
    public abstract T get(int i);  
    public abstract Backend<T> set(int i, T x);  
}
```

get an element for a path

set an element for a path

28 methods for free based on these, e.g.:

has, forEach, forEachNonNull,
nonNullIndices, findFirst, findLast,
findAny, clearAll, ...



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

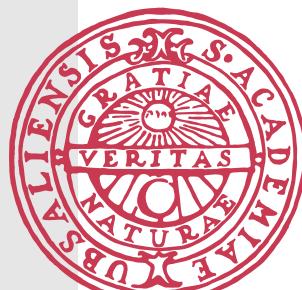
Fast access and insert, sparse

Array storage — uses oath ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

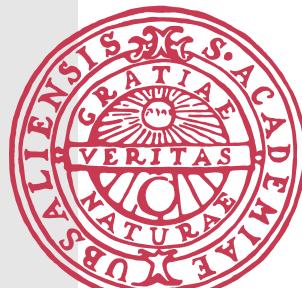
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

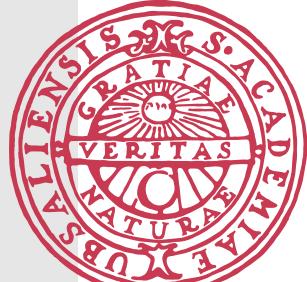
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

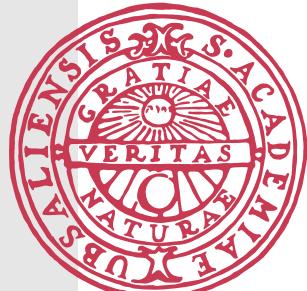
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

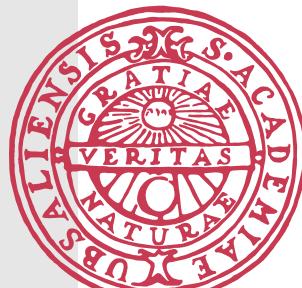
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

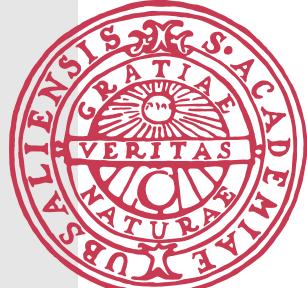
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

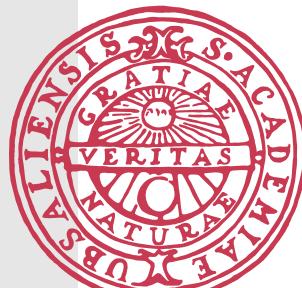
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Backend Implementations

- Manually implemented, following a Backend interface with 2 abstract methods and 28 default implementations

Hash table storage — using path ids as keys

Sparse

Immutable storage: using path ids as keys in an immutable map

Cheap copies

Indexed storage: like hash table storage, but with a backwards-map to permit fast location lookups of *values*

Fast way to find position by hash code

Sorted array storage: one array of non-null indices pointing into second array of values

Sparse array with efficient in-order iteration

Trie storage — uses a bitwise trie, four levels deep, using path ids as keys

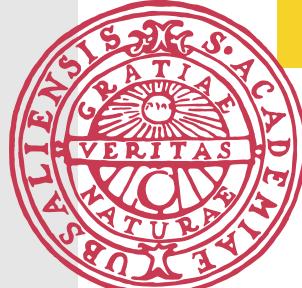
Fast access and insert, sparse

Array storage — uses path ids as array indices

Fast access and insert, dense

Reverse array storage — uses path ids as array indices

Like Array but reverse the data order



Frontend Implementations

- Three implementations for our case study: Sequence, Map and Matrix

Sequence: implements all methods in Java's List interface with ~4 SLOC per method

Map: implementation of AbstractMap (25 methods) in 87 SLOC

Matrix: see paper for details – no standard Java matrix exist

💀 Anecdotal evidence! 💀

C_b can be used to implement full-featured collections

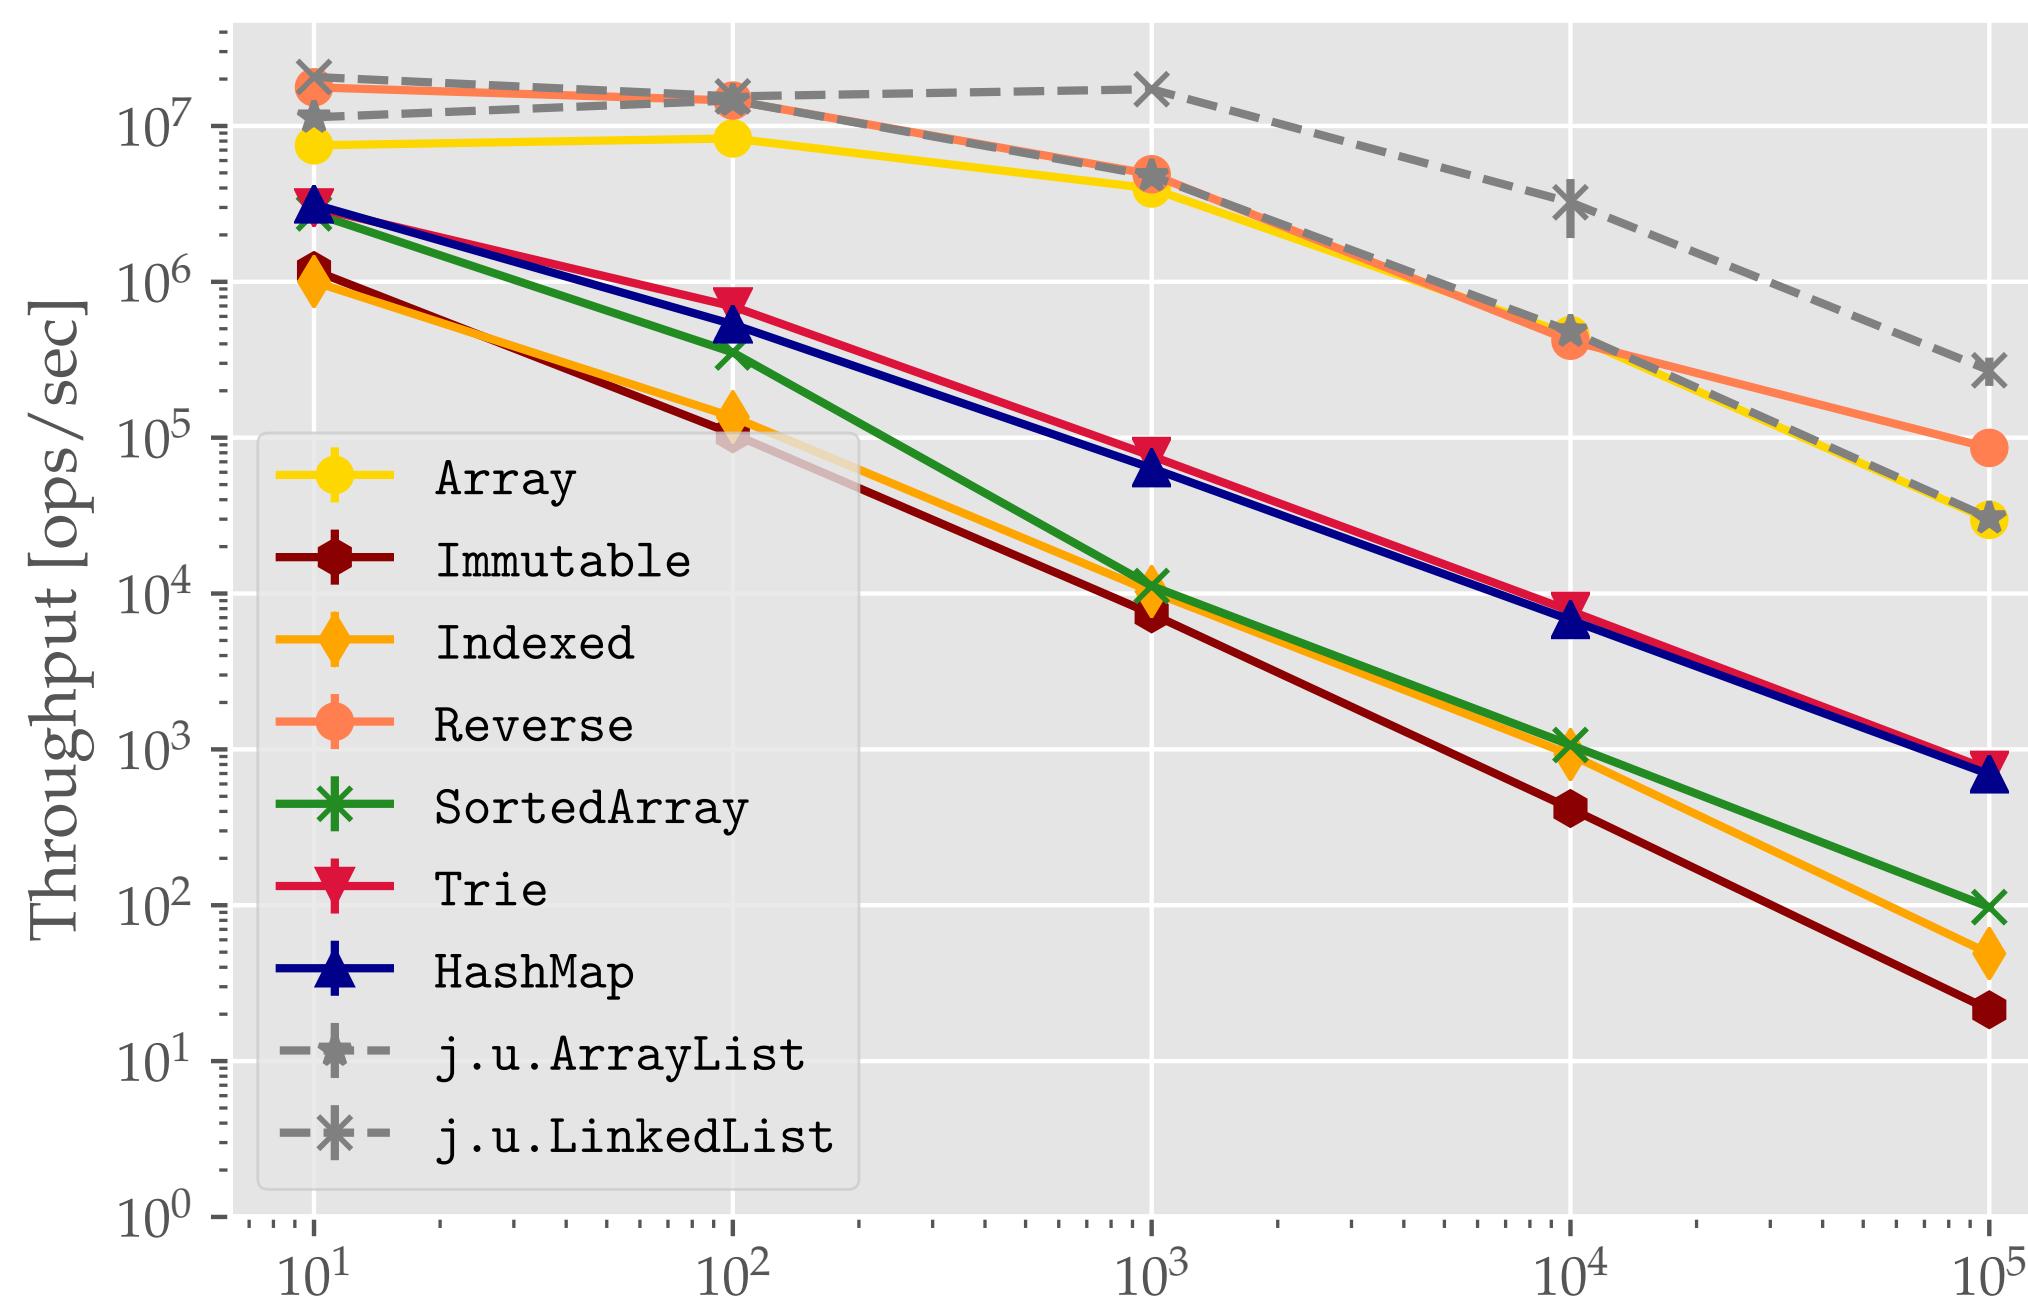
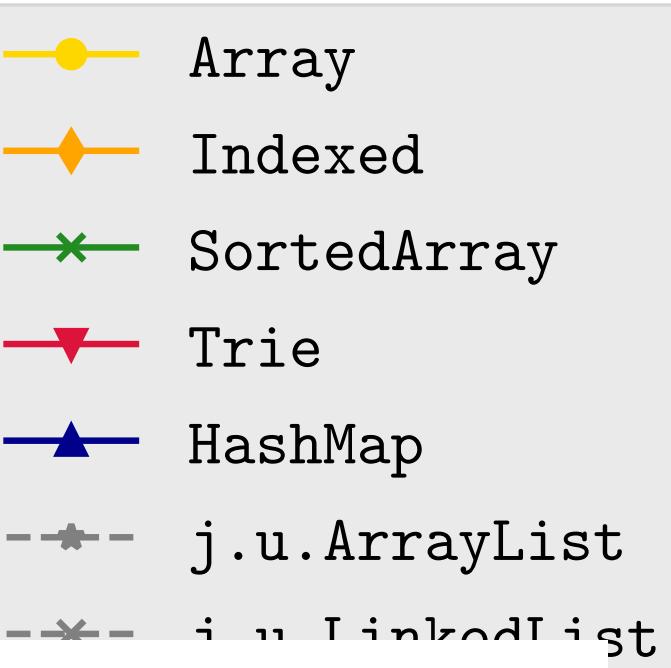
C_b implementations are reasonably straightforward

	Frontends		
Backends	Seq	Map	Matrix
Hashmap	👍	👍	👍
Immutable	👍	👍	👍
Indexed	👍	👍	👍
SortedArr	👍	👍	👍
Trie	👍	👍	👍
Array	👍	👍	👍
Reverse	👍	👍	👍

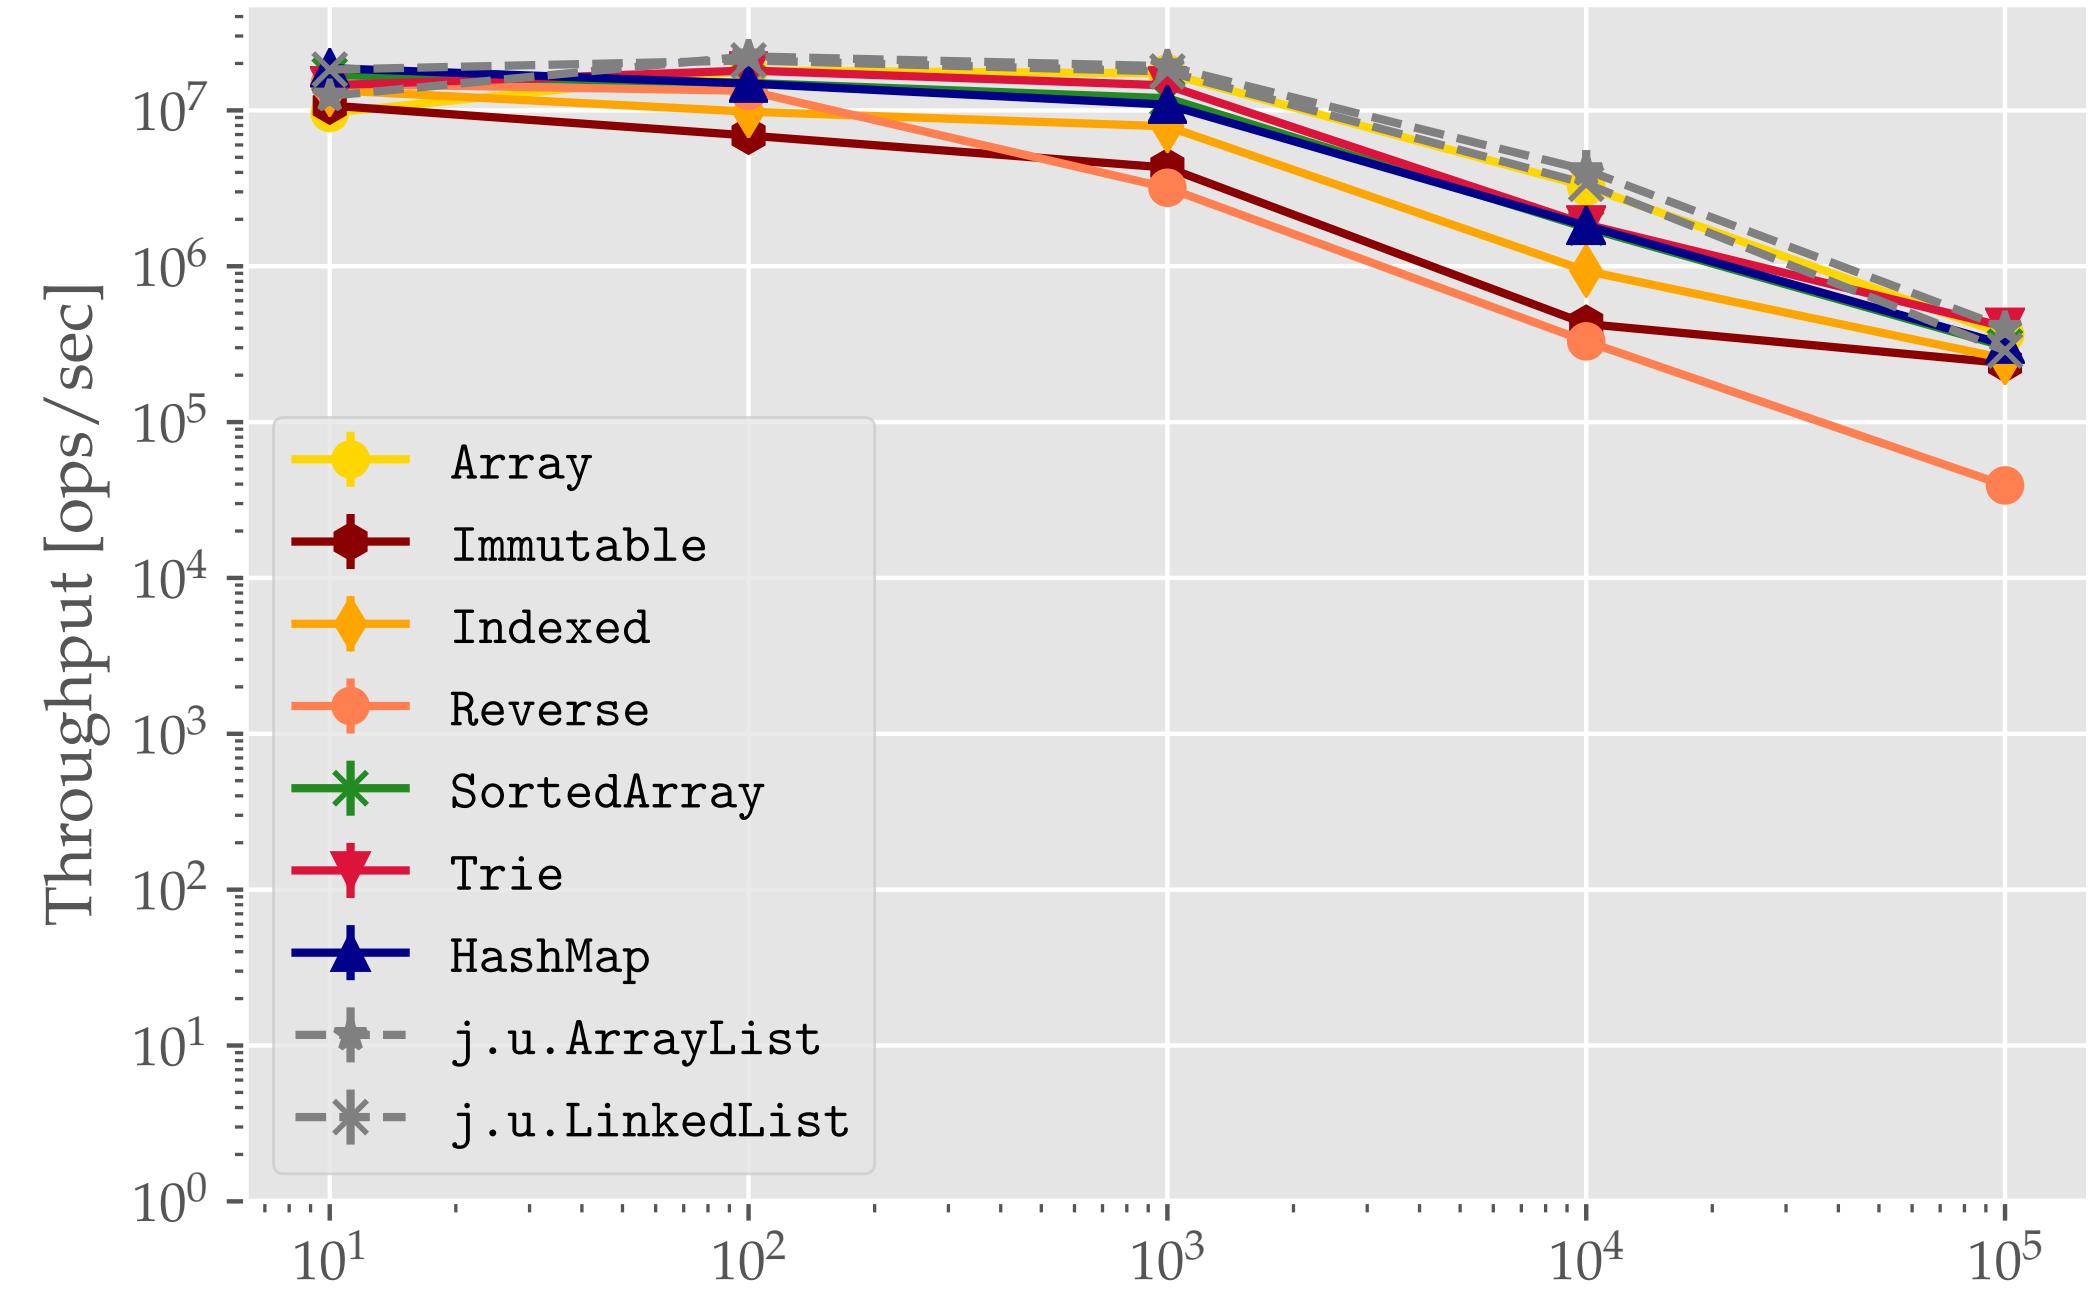
Memory-inefficient



Performance Compared to Java Collections

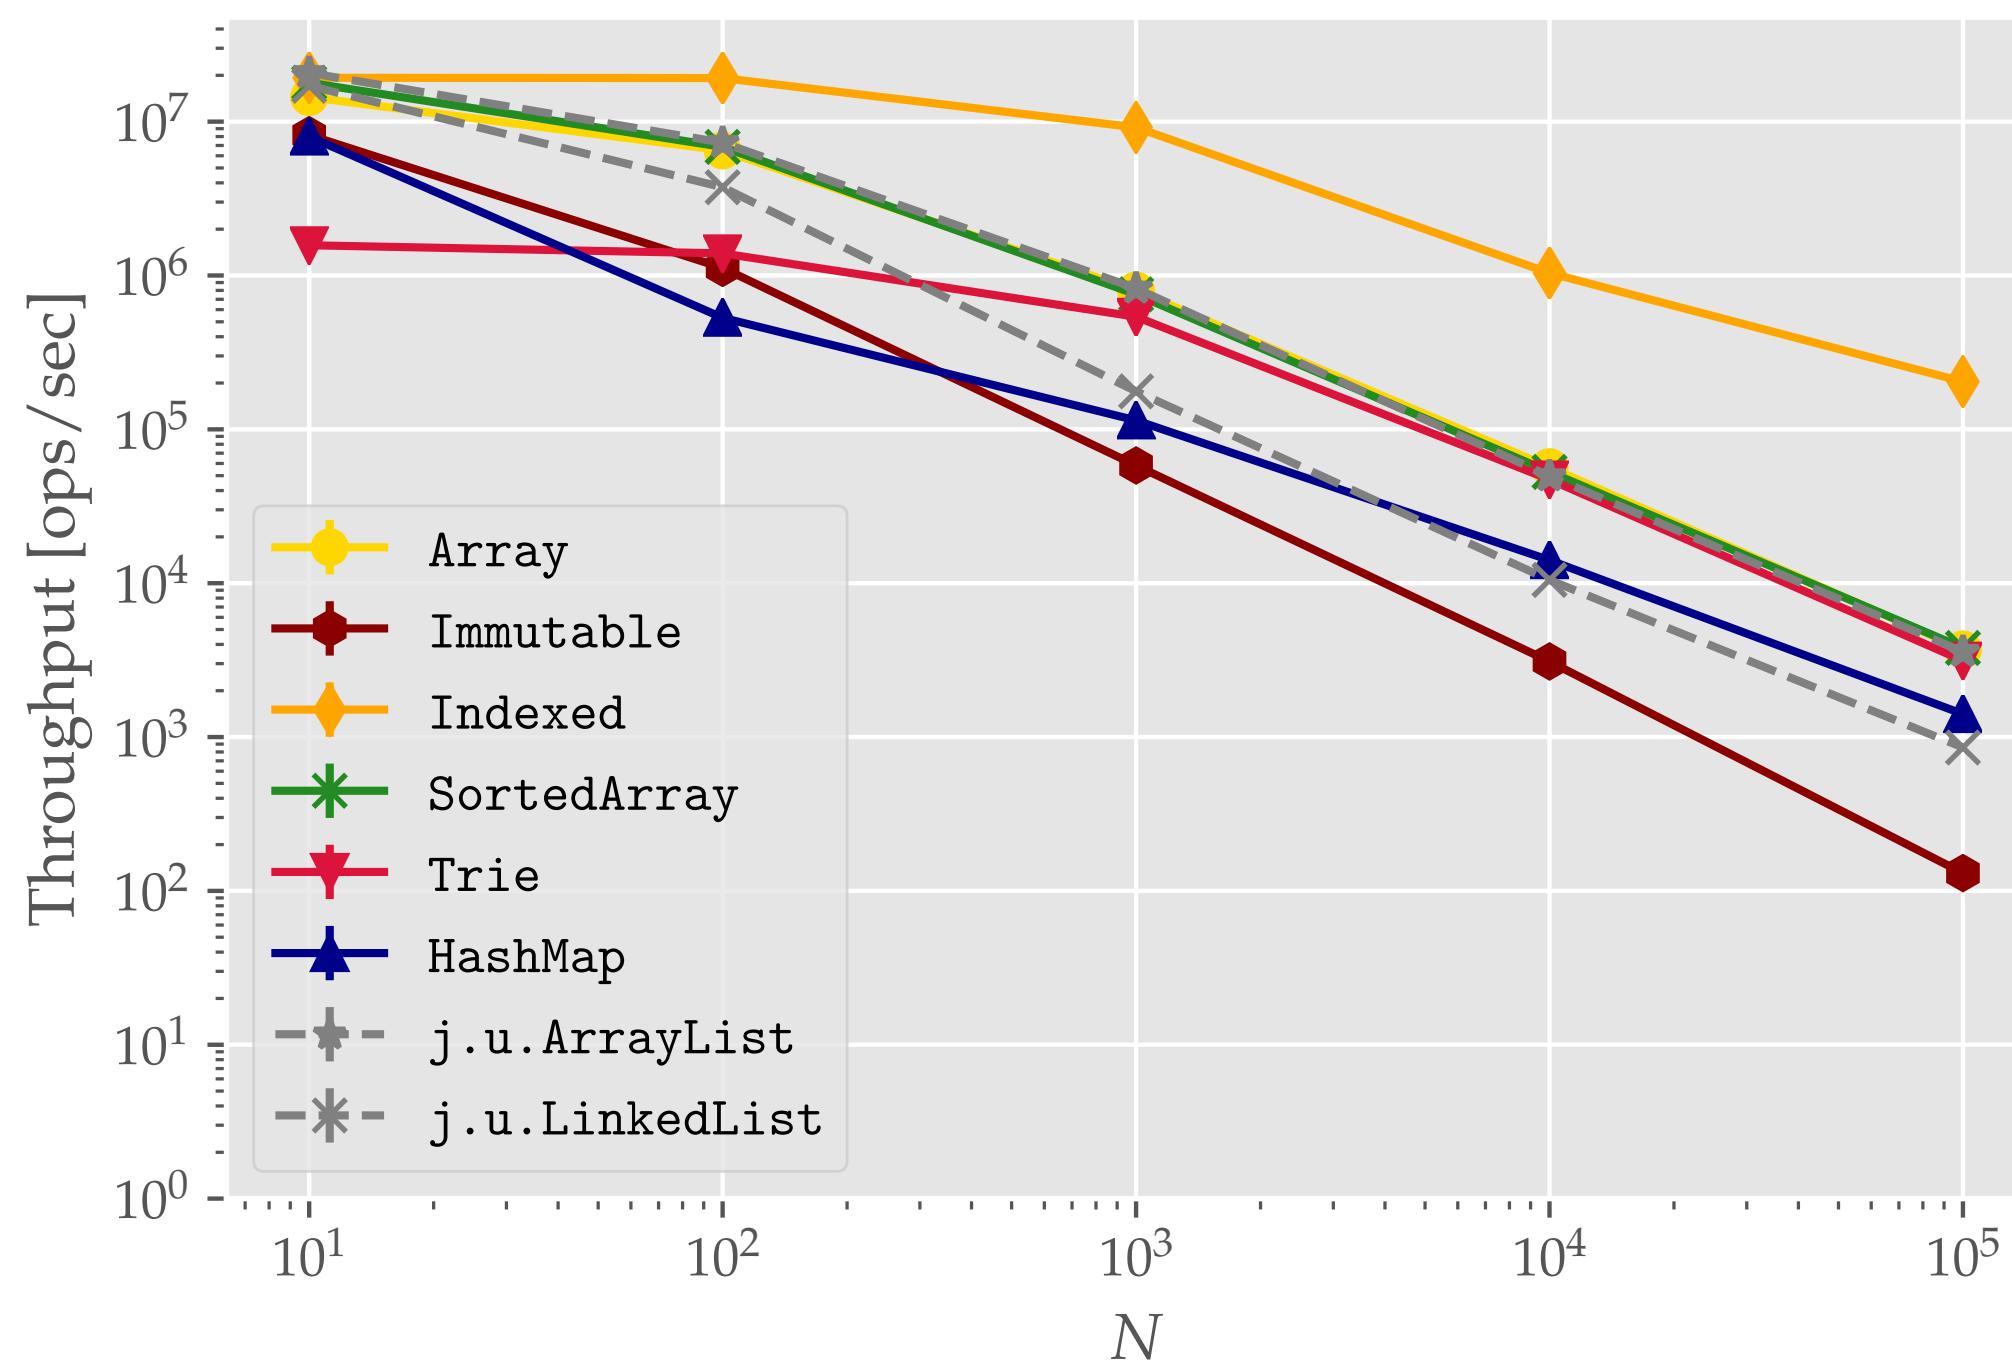
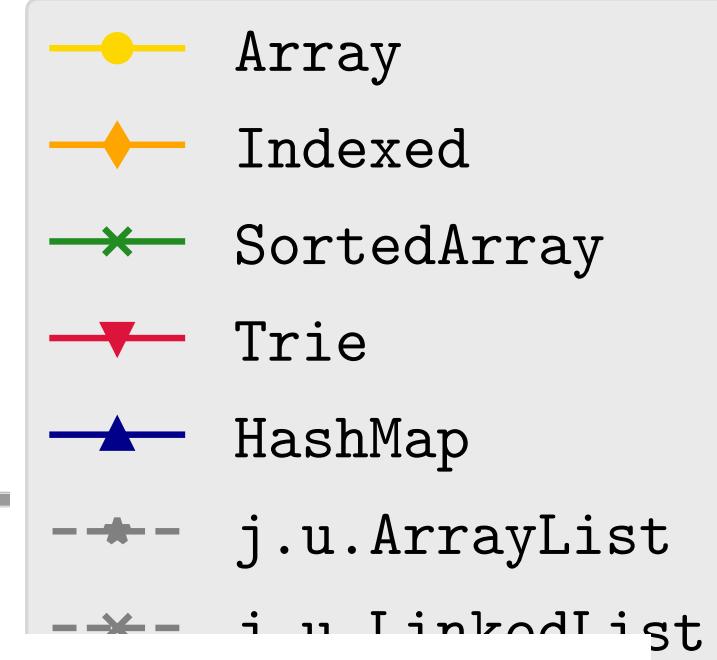


Insertion at front of a list

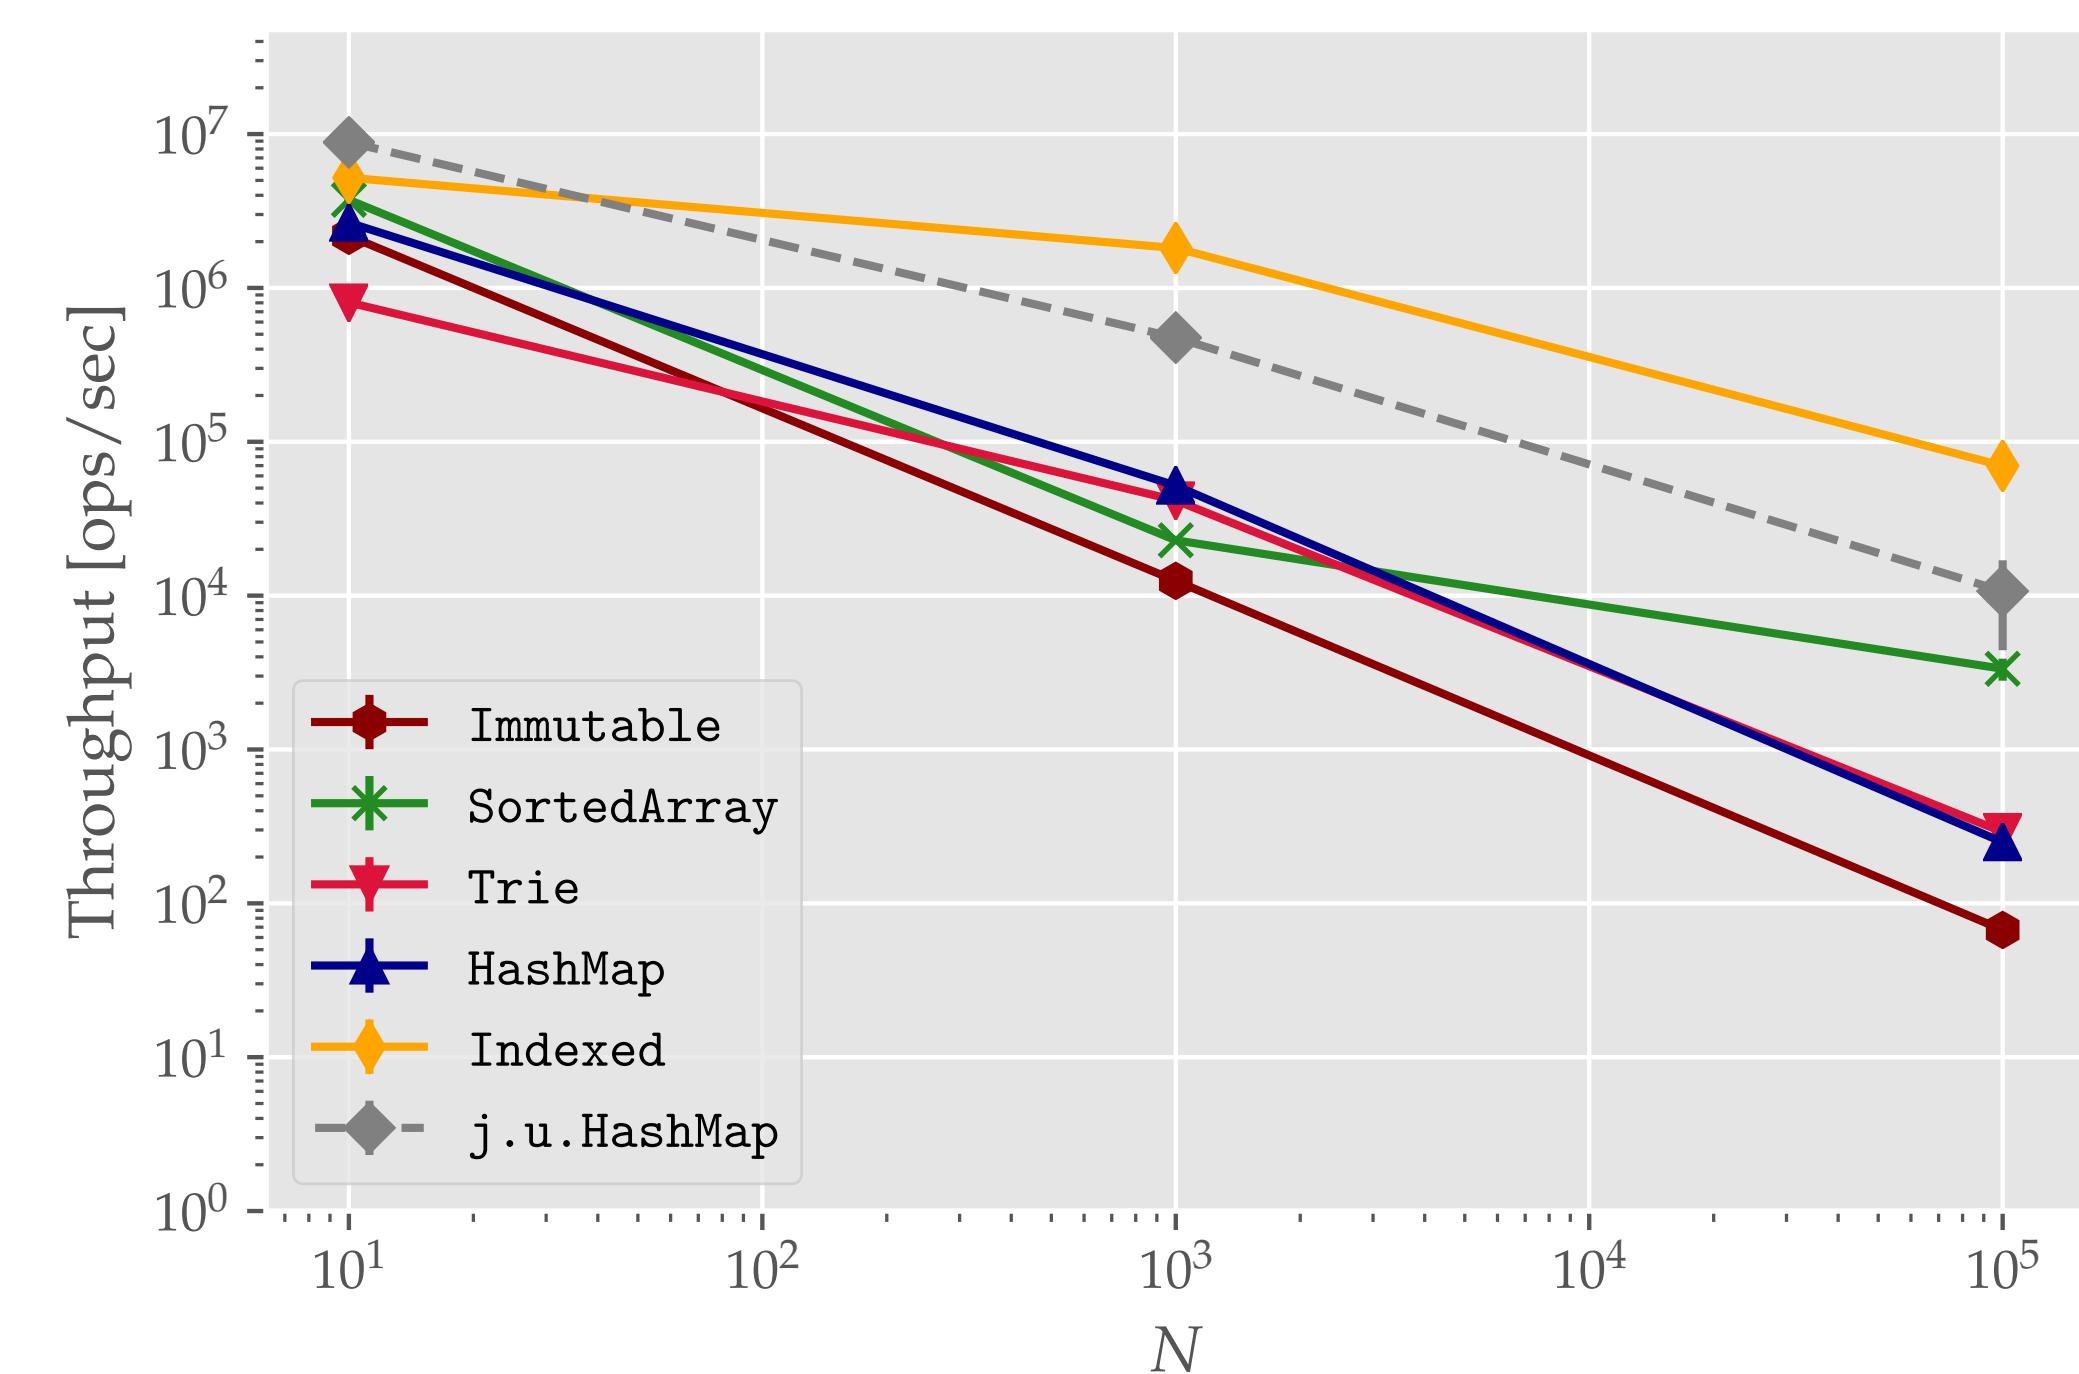


Insertion at end of a list

Performance Compared to Java Collections

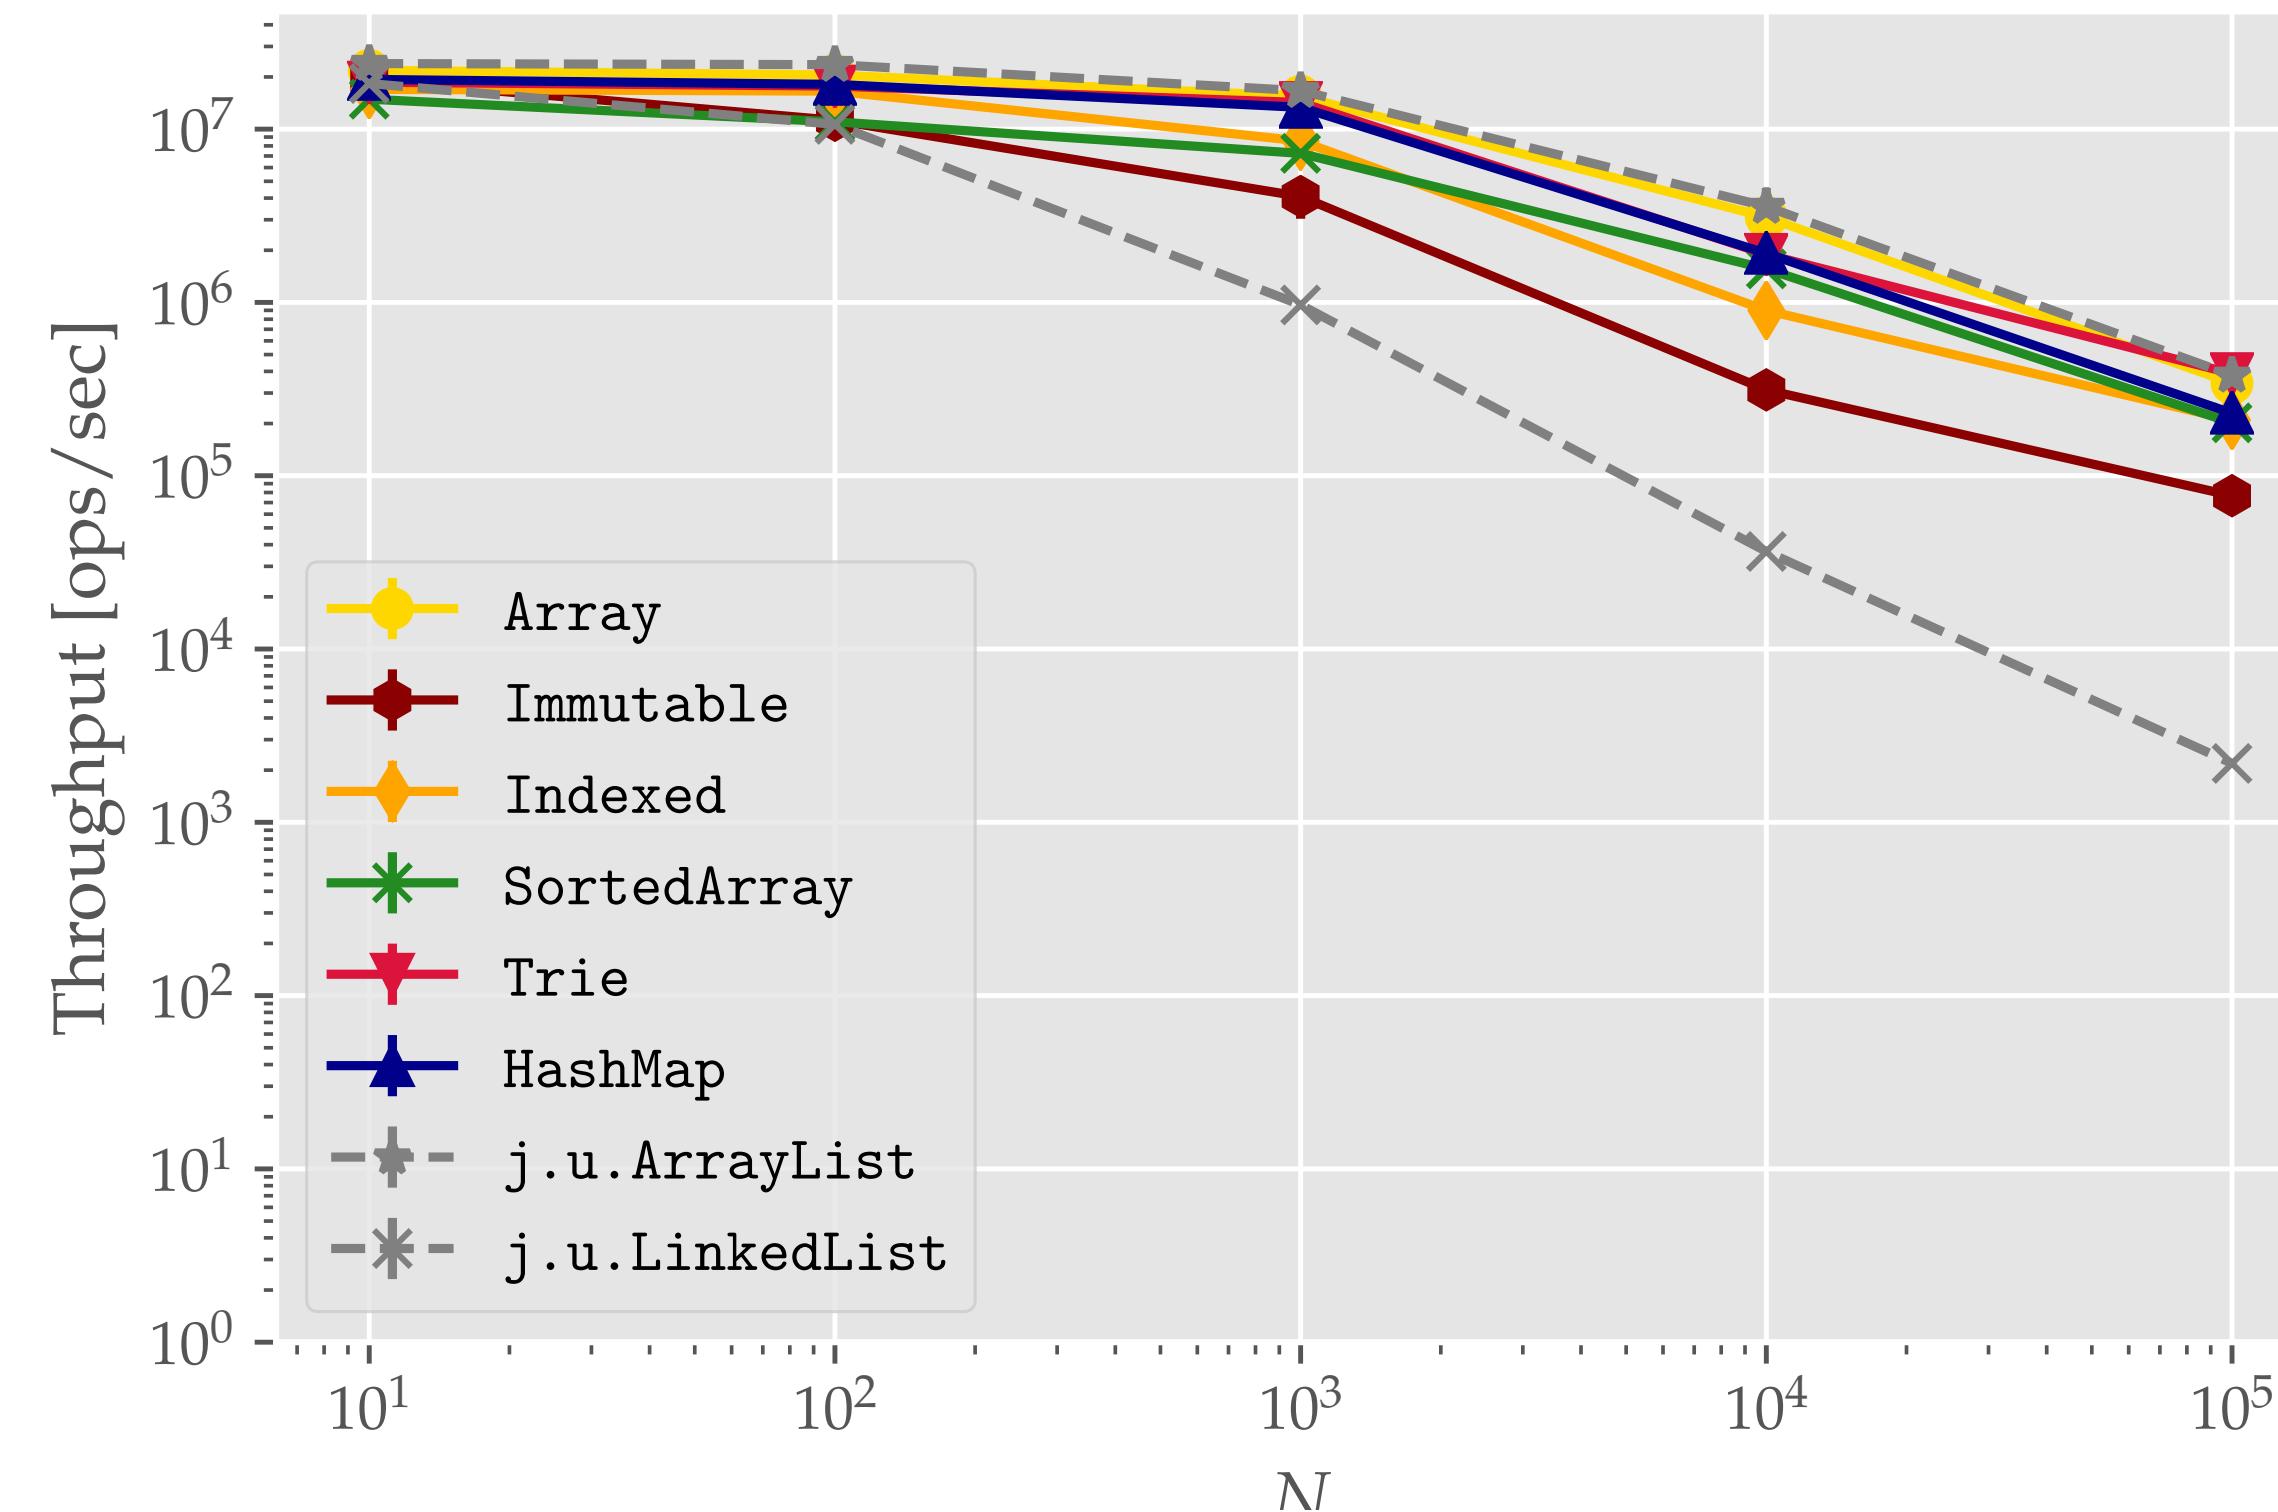


Searching for non-existing value in a list

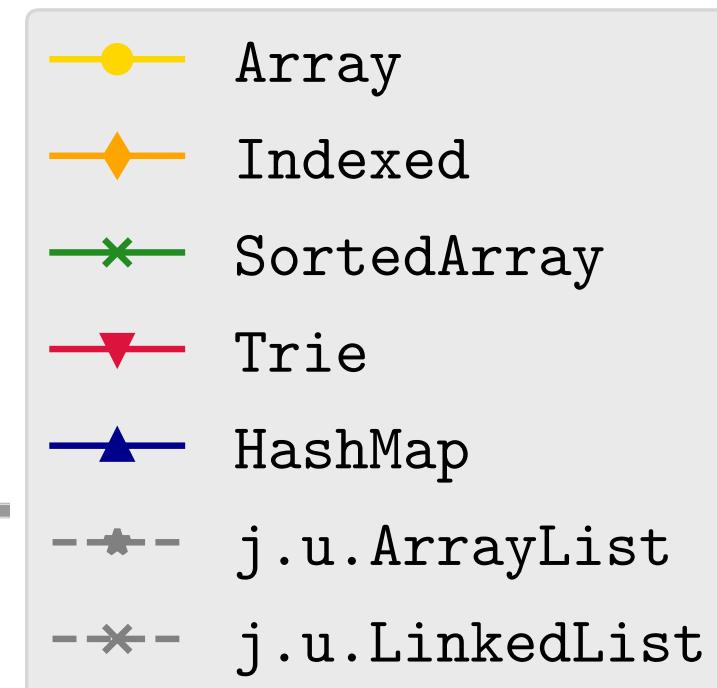


Searching for a non-existing value in a map

Performance Compared to Java Collections

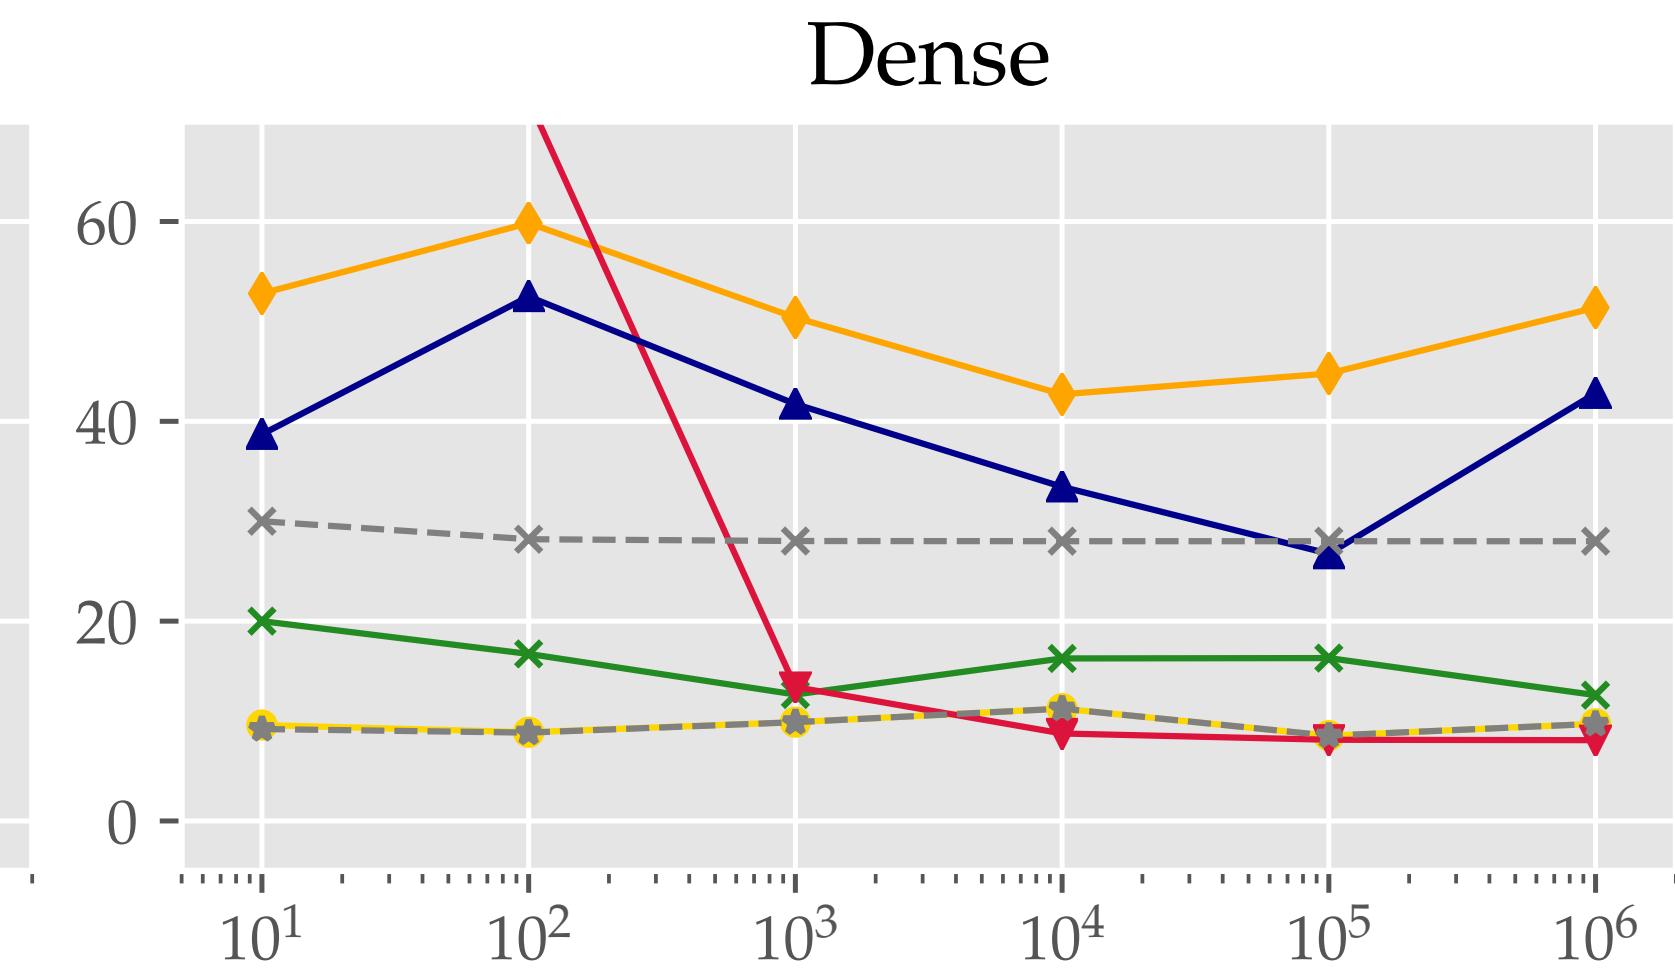
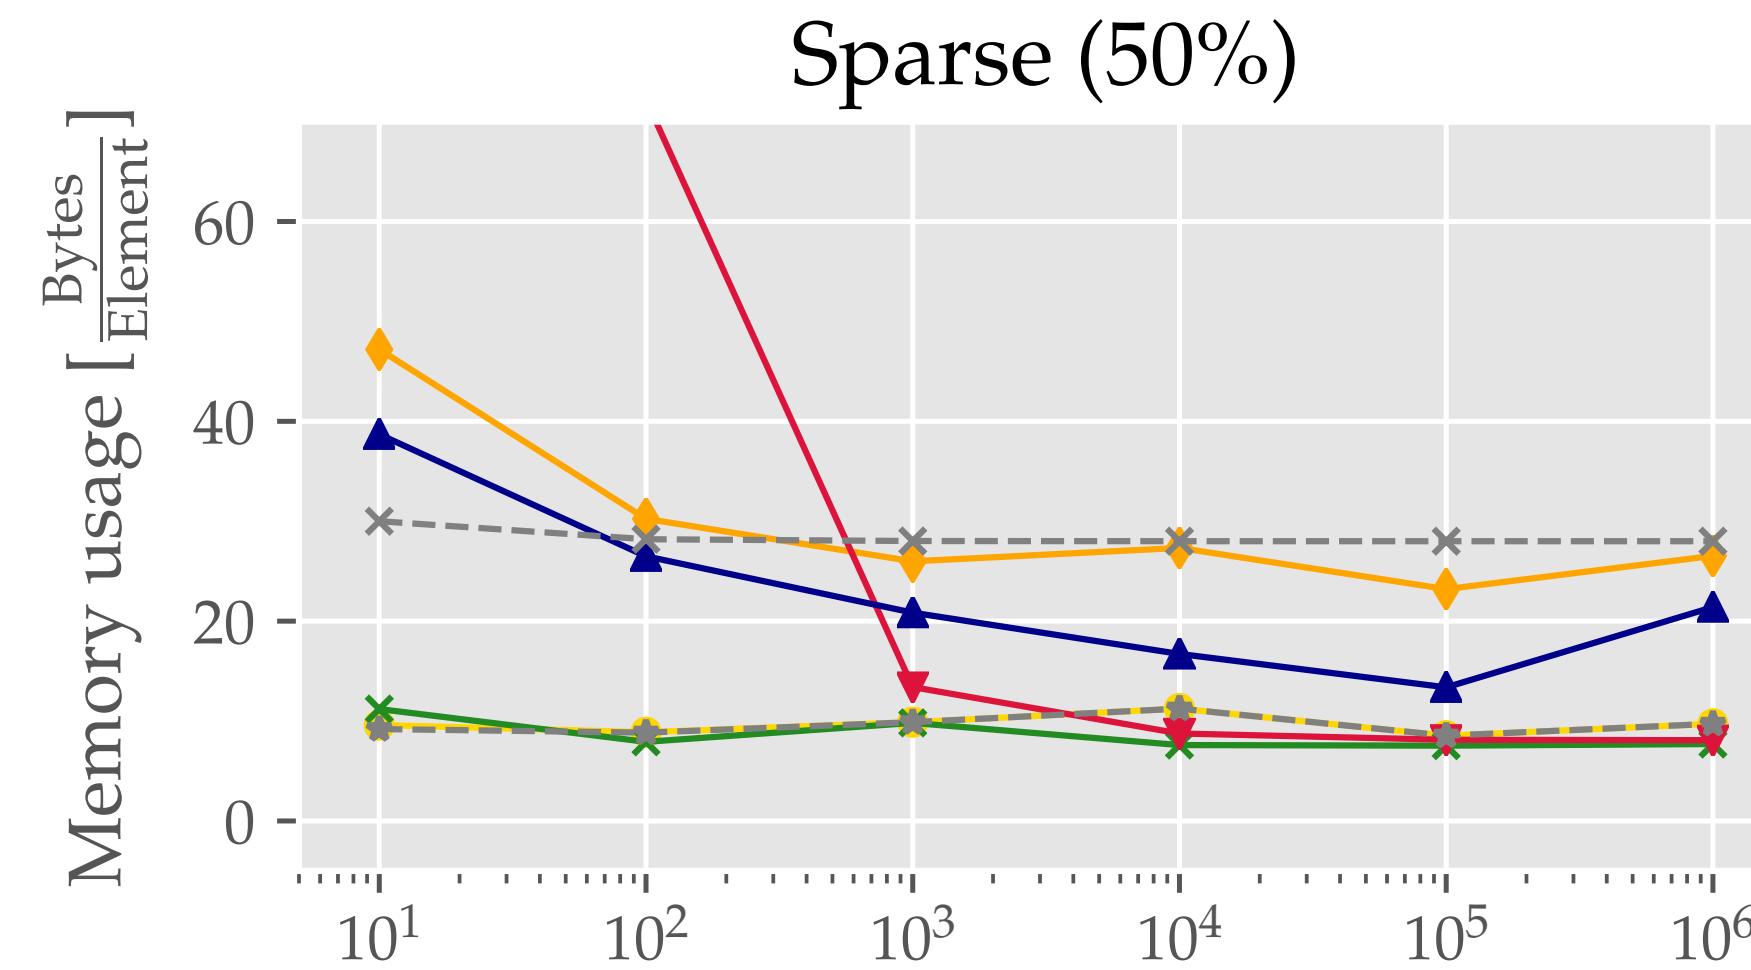
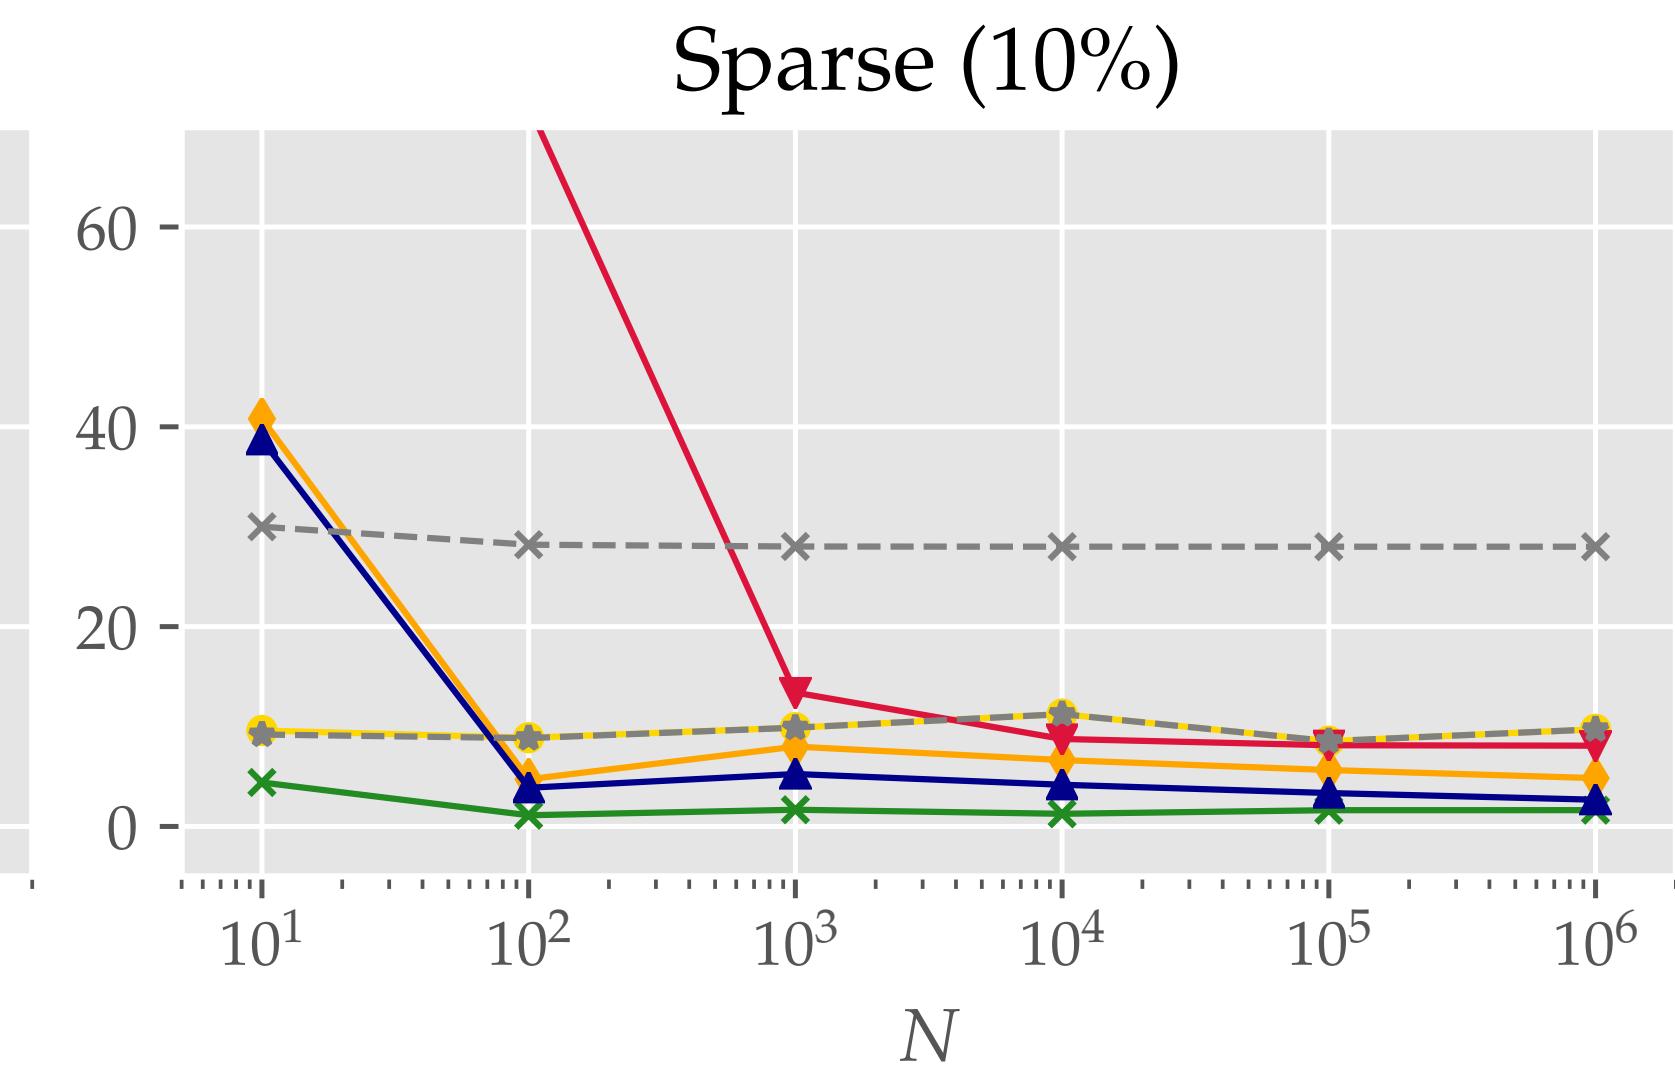
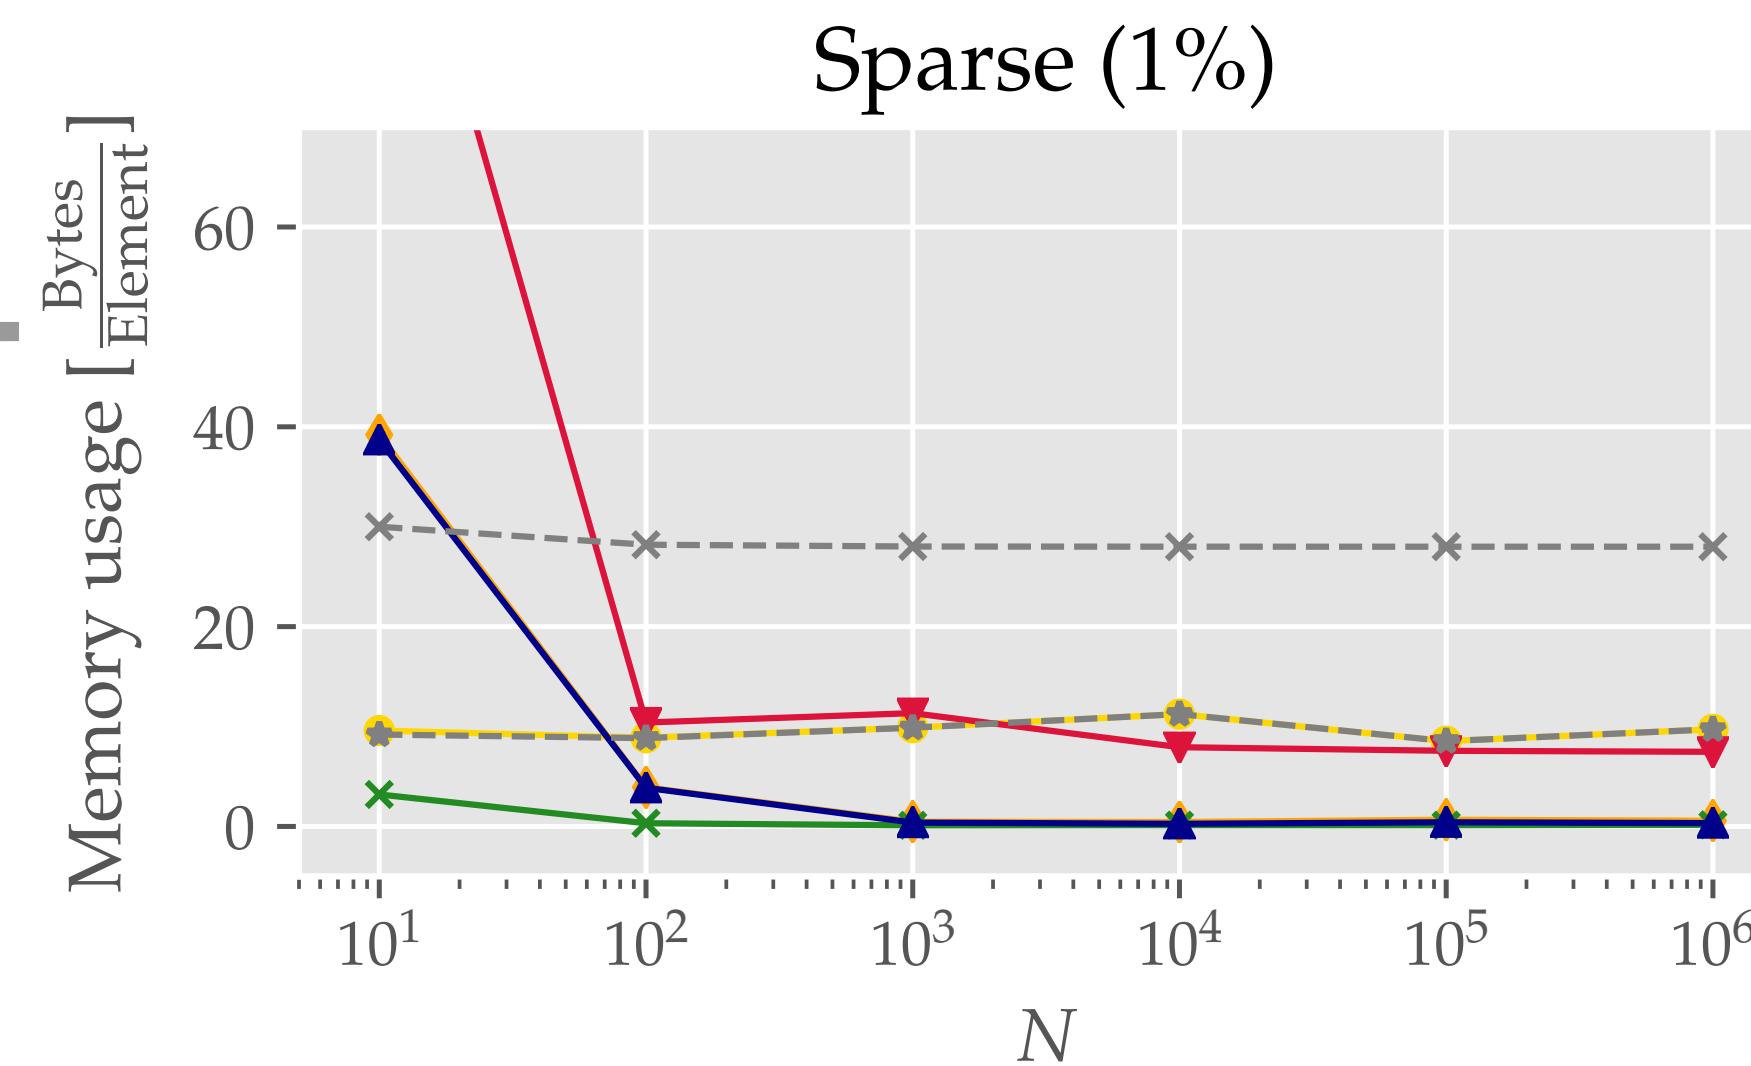
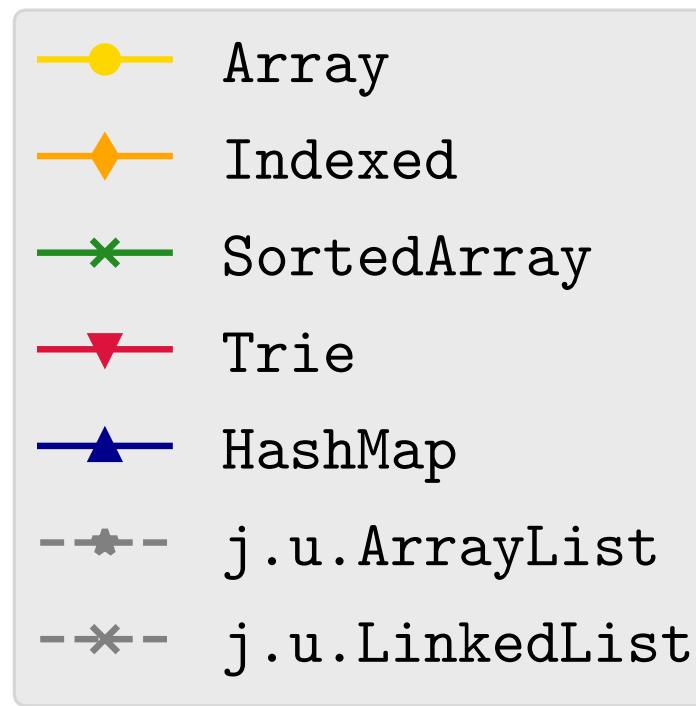


Random Access



Even though we insert an extra level of indirection for operations, performance is still roughly on par with native Java implementations.
(Naturally not for all backends!)

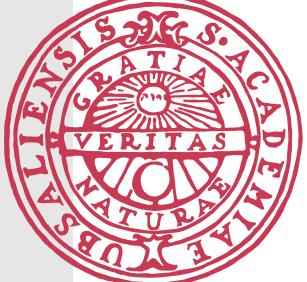
Also in the Paper: Tunability



Sequence sparseness vs. length vs. memory usage.

In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation

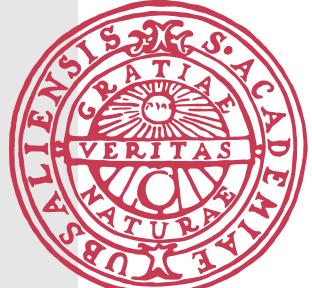


In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation
- Frontends and backends can be combined freely

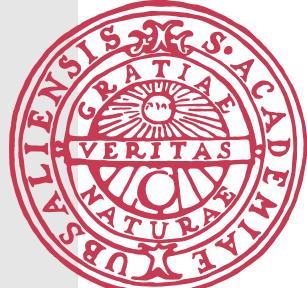
Allows backend optimizations to be reused across many collections

Allows tuning collections and reusing new backends with existing frontends



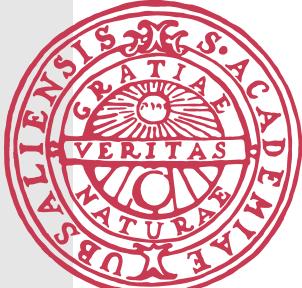
In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation
- Frontends and backends can be combined freely
 - Allows backend optimizations to be reused across many collections
 - Allows tuning collections and reusing new backends with existing frontends
- A simple automaton that controls how the frontend may navigate the space of the backend allows us to automatically generate the code for connecting any frontend with any backend



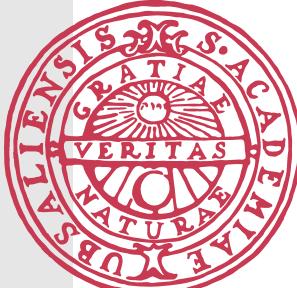
In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation
- Frontends and backends can be combined freely
 - Allows backend optimizations to be reused across many collections
 - Allows tuning collections and reusing new backends with existing frontends
- A simple automaton that controls how the frontend may navigate the space of the backend allows us to automatically generate the code for connecting any frontend with any backend
- We believe the code is straightforward, although it is initially unnatural if you are used to implementing data structures in the usual style



In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation
- Frontends and backends can be combined freely
 - Allows backend optimizations to be reused across many collections
 - Allows tuning collections and reusing new backends with existing frontends
- A simple automaton that controls how the frontend may navigate the space of the backend allows us to automatically generate the code for connecting any frontend with any backend
- We believe the code is straightforward, although it is initially unnatural if you are used to implementing data structures in the usual style
- Our simple design scales to implementing the full interfaces of standard Java collections (only 2 methods needed in a backend)



In Summary

- We have proposed a new way for implementing collection data structures that separates functional specification from data representation
- Frontends and backends can be combined freely
 - Allows backend optimizations to be reused across many collections
 - Allows tuning collections and reusing new backends with existing frontends
- A simple automaton that controls how the frontend may navigate the space of the backend allows us to automatically generate the code for connecting any frontend with any backend
- We believe the code is straightforward, although it is initially unnatural if you are used to implementing data structures in the usual style
- Our simple design scales to implementing the full interfaces of standard Java collections (only 2 methods needed in a backend)

Thank you!

