



Tracing as a Service

stephan.brandauer@it.uu.se *

tobias.wrigstad@it.uu.se




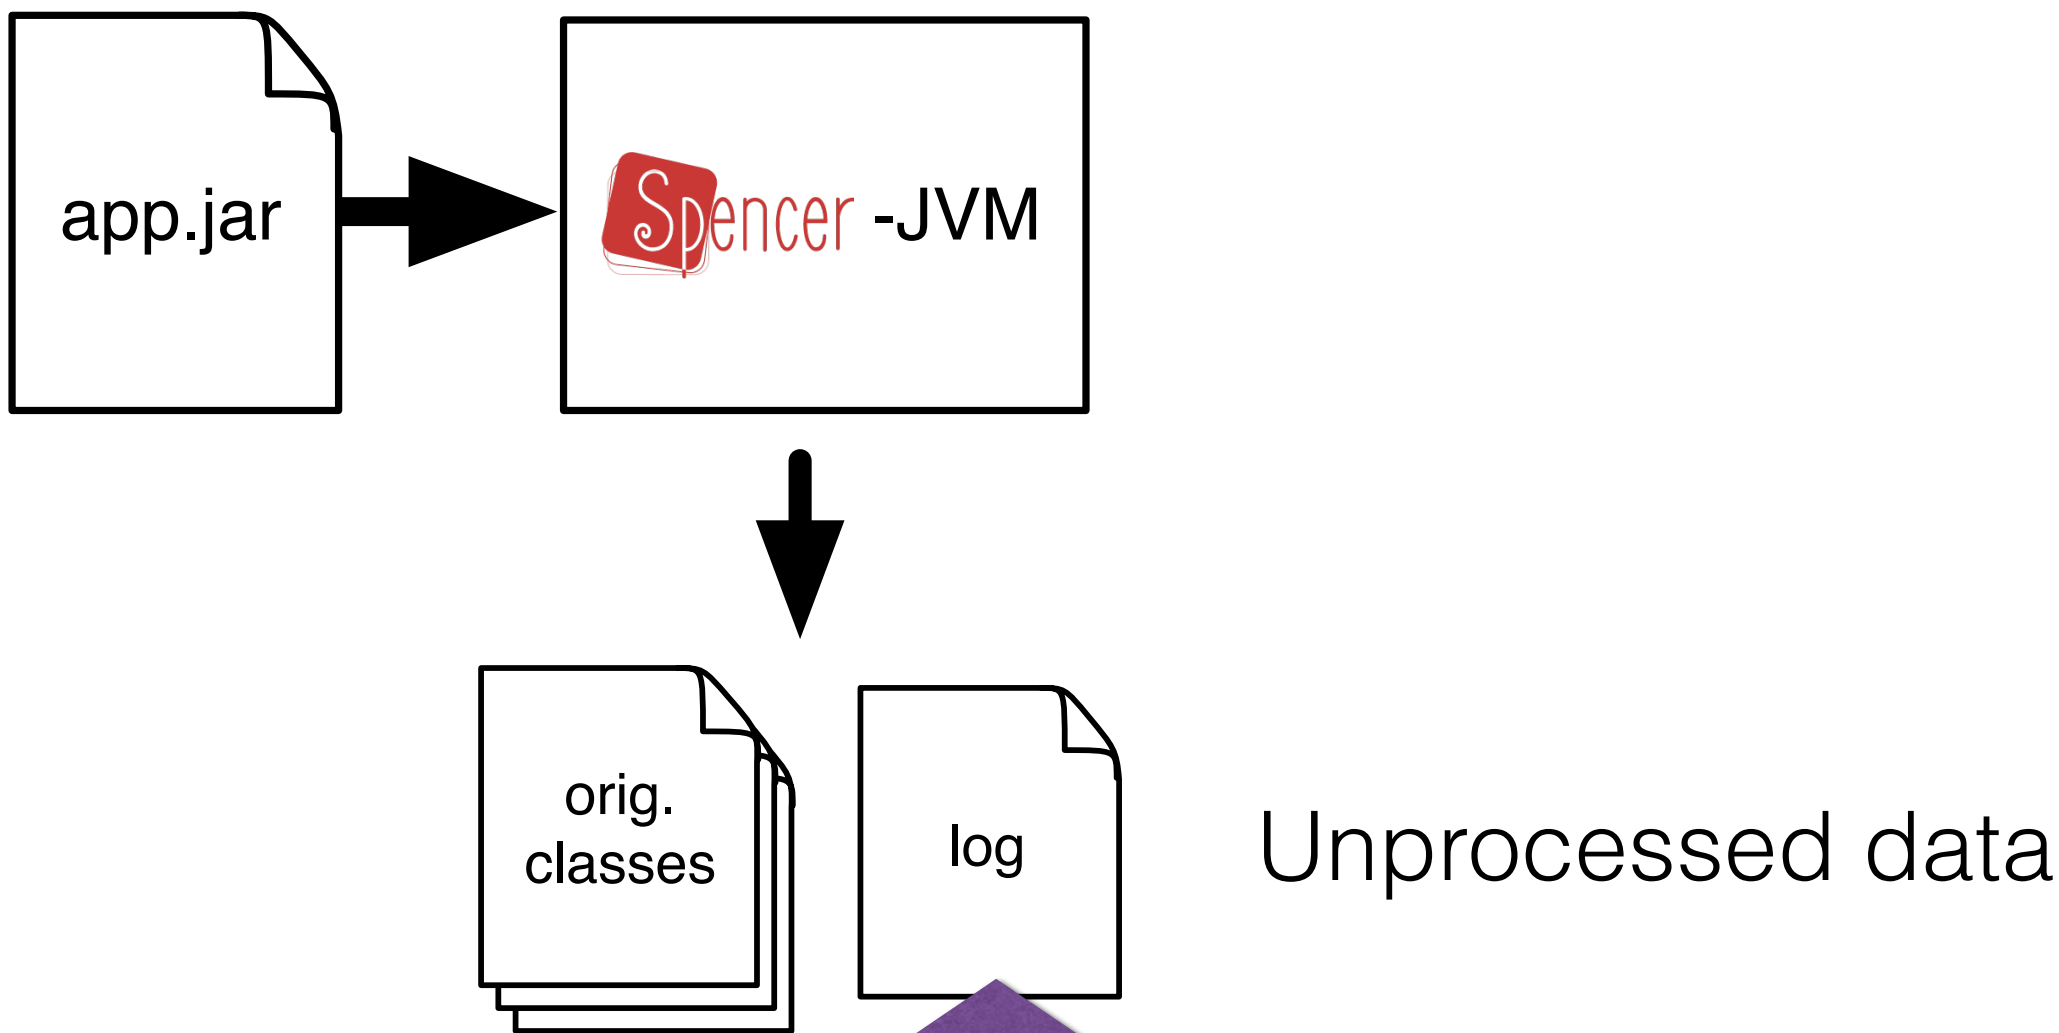
Web based service to query program traces.

“What do typical programs look like?”



Java bytecode

Workflow







```
#511073: ==> (java/lang/String @ 10247) . startsWith(Ljava/lang/String;)Z , callsite=MetaIndex.java:242 , thread=main
#511074: varstore - caller=java/lang/String :: startsWith @ 10247 var 1 , value was 0 , now is 10452 , thread=main
#511075: varload - caller=java/lang/String @ 10247 , var 1 , val=10452 , thread=main
#511076: ==> (java/lang/String @ 10247) . startsWith(Ljava/lang/String;I)Z , callsite=String.java:1434 , thread=main
#511077: varstore - caller=java/lang/String :: startsWith @ 10247 var 1 , value was 0 , now is 10452 , thread=main
#511078: fieldLoad - caller=java/lang/String :: startsWith @ 10247 , holder=java/lang/String @ 10247 , field=[C value , thread=main
#511079: varstore - caller=java/lang/String :: startsWith @ 10247 var 3 , value was 0 , now is 10248 , thread=main
#511080: varload - caller=java/lang/String @ 10247 , var 1 , val=10452 , thread=main
#511081: fieldLoad - caller=java/lang/String :: startsWith @ 10247 , holder=java/lang/String @ 10452 , field=[C value , thread=main
#511082: varstore - caller=java/lang/String :: startsWith @ 10247 var 5 , value was 0 , now is 10453 , thread=main
#511083: varload - caller=java/lang/String @ 10247 , var 1 , val=10452 , thread=main
#511084: fieldLoad - caller=java/lang/String :: startsWith @ 10247 , holder=java/lang/String @ 10452 , field=[C value , thread=main
#511085: fieldLoad - caller=java/lang/String :: startsWith @ 10247 , holder=java/lang/String @ 10247 , field=[C value , thread=main
#511086: varload - caller=java/lang/String @ 10247 , var 3 , val=10248 , thread=main
#511087: readmodify - callee=[C @ 10248 , caller=java/lang/String @ 10247 reads _0
#511088: varload - caller=java/lang/String @ 10247 , var 5 , val=10453 , thread=main
#511089: readmodify - callee=[C @ 10453 , caller=java/lang/String @ 10247 reads _0
#511090: <== (??? @ java/lang/String) . startsWith(???), thread=main
#511091: <== (??? @ java/lang/String) . startsWith(???), thread=main
```




Unprocessed data



preprocess+load




SQL-DB



calls ✓

```
# SELECT * FROM calls WHERE callstart = 511073 ;
```

caller	callee	name	callstart	callend	callsitefile	callsiteline	thread
10530	10247	startsWith	511073	511091	MetaIndex.java	242	main




calls ✓
uses ✓

```
# SELECT * FROM calls WHERE callstart = 511073 ;
```

caller	callee	name	callstart	callend	callsitefile	callsiteline	thread
10530	10247	startsWith	511073	511091	MetaIndex.java	242	main

```
# SELECT * FROM uses WHERE idx ≥ 511073 AND idx ≤ 511091 ;
```

caller	callee	name	method	kind	idx	thread
10247	10247	var_1	startsWith	varstore	511074	main
10247	10247	var_1	startsWith	varload	511075	main
... snip ...						
10247	10247	var_5	startsWith	varload	511088	main
10247	10453	_0	startsWith	read	511089	main



calls ✓
uses ✓
refs ✓

```
# SELECT * FROM calls WHERE callstart = 511073 ;
caller | callee | name      | callstart | callend | callsitefile | callsiteline | thread
-----+-----+-----+-----+-----+-----+-----+-----
10530  | 10247  | startsWith | 511073    | 511091  | MetaIndex.java | 242          | main
```

```
# SELECT * FROM uses WHERE idx ≥ 511073 AND idx ≤ 511091 ;
caller | callee | name  | method  | kind   | idx  | thread
-----+-----+-----+-----+-----+-----+-----
10247  | 10247  | var_1 | startsWith | varstore | 511074 | main
10247  | 10247  | var_1 | startsWith | varload  | 511075 | main
... snip ...
10247  | 10247  | var_5 | startsWith | varload  | 511088 | main
10247  | 10453  | _0    | startsWith | read     | 511089 | main
```

```
# SELECT * FROM refs WHERE caller = 10247 AND kind = 'field' ;
caller | callee | kind  | name  | refstart | refend | thread
-----+-----+-----+-----+-----+-----+-----
10247  | 10248  | field | value | 421877  |        | main
```

Queries

- Spencer uses a query DSL.
 - Compiled to SQL, and cached.
 - Makes caching effective.
 - Easier to use.
 - No **NEED** for SQL — but better performance.
- PostgreSQL is surprisingly expressive!

ImmutableObj()

```
SELECT id FROM objects WHERE id > 4
EXCEPT
  (SELECT DISTINCT callee AS id
   FROM uses_cstore
   WHERE callee > 4
   AND NOT(caller = callee AND method = '<init>')
   AND (kind = 'fieldstore' OR kind = 'modify'))
```

HeapDeeply(ImmutableObj())

```
(
  SELECT id FROM objects WHERE id > 4
EXCEPT
  (SELECT DISTINCT callee AS id
FROM uses_cstore
WHERE
  callee > 4 AND
  NOT(caller = callee AND method = '<init>') AND
  (kind = 'fieldstore' OR kind = 'modify'))

) INTERSECT (
  SELECT id FROM objects WHERE id > 4
EXCEPT
  (WITH RECURSIVE canheapreach(id) AS (
    SELECT id FROM objects WHERE id > 4
  EXCEPT
    (SELECT id FROM objects WHERE id > 4
  EXCEPT
    ...
```

```
...
  (SELECT DISTINCT callee AS id
FROM uses_cstore
WHERE
  callee > 4 AND
  NOT(caller = callee AND method = '<init>') AND
  (kind = 'fieldstore' OR kind = 'modify'))
  )


UNION
  SELECT
    refs.caller AS id
  FROM refs
  JOIN canheapreach ON canheapreach.id = refs.callee
  WHERE kind = 'field'
)
SELECT id FROM canheapreach)

)
```


API + Meta Info

- There exists an API that gives you results (formatted as JSON objects)
- <http://www.spencer-t.racing/doc/api>
- Meta info: “per-object tags”
 - class
 - allocation site + “time”
 - coming: number of reads, writes
 - I’m happy to add more

Example from API



(g) This shows, for each property, the average across data sets of the proportion of classes that *only* (green, right)/*never* (red, left) produced instances that had the property. We infer, unsoundly, that there exists an invariant that guarantees the property holds statically. Classes with less than 10 instances are ignored.

Status

- Tracing tool implementation: done, modulo maintenance
- Analysis DSL: useful for some use cases. New uses through extension.
- Web interface: work in progress, has performance issues

Sources

- <http://spencer-t.racing>
- “Spencer: Interactive Heap Analysis for the Masses”, to appear (International Conference on Mining Software Repositories)
- “Mining for Safety using Interactive Trace Analysis”, to appear (Workshop on Quantitative Aspects of Programming Languages and Systems)

Questions?